

US010650544B2

(12) United States Patent

Young et al.

(54) OPTIMIZED SHADOWS IN A FOVEATED RENDERING SYSTEM

(71) Applicant: Sony Interactive Entertainment Inc.,

Tokyo (JP)

(72) Inventors: **Andrew Young**, San Mateo, CA (US);

Chris Ho, San Mateo, CA (US); Jeffrey Roger Stafford, San Mateo, CA

(US)

(73) Assignee: Sony Interactive Entertainment Inc.,

Tokyo (JP)

(*) Notice: Subject to any disclaimer, the term of this

patent is extended or adjusted under 35

U.S.C. 154(b) by 0 days.

(21) Appl. No.: 15/727,505

(22) Filed: Oct. 6, 2017

(65) Prior Publication Data

US 2018/0357780 A1 Dec. 13, 2018

Related U.S. Application Data

- (60) Provisional application No. 62/517,829, filed on Jun. 9, 2017.
- (51) Int. Cl. G06K 9/00 (2006.01) G06T 7/66 (2017.01)

(Continued)

(Continued)

(58) Field of Classification Search

CPC G06T 7/66; G06T 7/20; G06T 7/0002; G06T 7/73; G06T 11/20; G06T 13/80 See application file for complete search history.

(10) Patent No.: US 10,650,544 B2

(45) **Date of Patent:** May 12, 2020

(56) References Cited

U.S. PATENT DOCUMENTS

OTHER PUBLICATIONS

Patney et (Towards Foveated Rendering for Gaze-Tracked Virtual Reality, ACM Trans Graphic vol. 35, No. 6, Article 176, Nov. 2016, p. 1-12), (Year: 2016).*

(Continued)

Primary Examiner — Amara Abdi (74) Attorney, Agent, or Firm — Penilla IP, APC

(57) ABSTRACT

A method for implementing a graphics pipeline. The method includes building a first shadow map of high resolution, and building a second shadow map based on the first shadow map of lower resolution. The method includes determining a light source affecting a virtual scene, and projecting geometries of objects of an image of the virtual scene onto a plurality of pixels of a display from a first point-of-view. The method includes determining a foveal region when rendering the image, wherein the foveal region corresponds to where an attention of a user is directed. The method includes determining a first set of geometries is drawn to a first pixel, determining the first set of geometries is in shadow based on the light source, and determining the first set of geometries is outside of the foveal region. The method includes rendering the first set of geometries using the second shadow map.

18 Claims, 13 Drawing Sheets

Building a first shadow map of high resolution

Y

Building a second shadow map based on the first shadow map of low resolution

Y

Determining a light source affecting a virtual scene

Frojecting geometries of objects of an image of the virtual scene onto a plurality of pixels of a display from a first point of-view

Determining a first set of geometries is drawn to a first pixel

Determining the first set of geometries is in shadow based on the light source

Y

Determining a first set of geometries is in shadow based on the light source

Y

Determining the first set of geometries is in shadow based on the light source

Y

Determining the first set of geometries is in shadow based on the light source

Y

Determining the first set of geometries is outside of the foveal region

FRondering the first set of geometries is outside of the foveal region

(51)	Int. Cl.				
	G06T 7/00	(2017.01)			
	G06T 13/80	(2011.01)			
	G06T 7/73	(2017.01)			
	G06T 11/20	(2006.01)			
	G06T 7/20	(2017.01)			
	G06F 3/01	(2006.01)			
	G06T 11/40	(2006.01)			
	G06T 15/10	(2011.01)			
(52)	U.S. Cl.				
. /	CPC	G06T 7/20 (2013.01); G06T 7/73			
	(2017.01); G06T 11/20 (2013.01); G06T 11/40				
	206T 13/80 (2013.01); G06T 15/10				
	//); G06T 2210/36 (2013.01); G06T			
	(2215/12 (2013.01)			

(56) References Cited

U.S. PATENT DOCUMENTS

2016/0379381	A1*	12/2016	Krutsch G06T 11/001
			345/426
2016/0381050	A1*	12/2016	Shanbhogue G06F 9/30145
			726/23
2017/0263046	A1*	9/2017	Patney G06T 15/20
2018/0164592	A1*	6/2018	Lopes A61B 3/113
2018/0286105	A1*	10/2018	Surti H04N 5/3696

OTHER PUBLICATIONS

Weier et al ("Foveated Real-Time Ray Tracing for Head-Mounted Displays", Pacific Graphics 2016 E. Grinspun, B. Bickel, and Y. Dobashi) (Year: 2016).*

Stengel et al ("Adaptive Image-Space Sampling for Gaze-Contingent Real-time Rendering", Eurographics Symposium on Rendering 2016 E. Eisemann and E. Fiume (Year: 2016).*

Guenter, Brian et al: "Foveated 3D graphics", ACM Transactions on Graphics, vol. 31, No. 6, Nov. 1, 2012, 10 pages.

Patney, Anjul et al: "Towards foveated rendering for gaze-tracked virtual reality", ACM Transactions on Graphics (TOG) ACM, US, vol. 35, No. 6, Nov. 11, 2016, 12 pages.

Lauritzen, Andrew et al: "Layered Variance Shadow Maps", Graphics Interface 2005: Proceedings; Victoria, British Columbia, May 9-11, 2005, Canadian Information Processing Society, 403 King Street West, Suite 205 Toronto, Ont. M5U 1LS Canada, May 28, 2008, pp. 139-146.

Borza, Andrei, et al: "Eye Tracked Shadow Maps", 2014 IEEE 3rd Global Conference on Consumer Electronics (GCCE), IEEE, Oct. 7, 2014, pp. 757-761.

Stengel, Michael et al: "Adaptive Image-Space Sampling for Gaze-Contingent Real-time Rendering", Computer Graphics Forum, vol. 35, No. 4, Jul. 1, 2016, pp. 129-139.

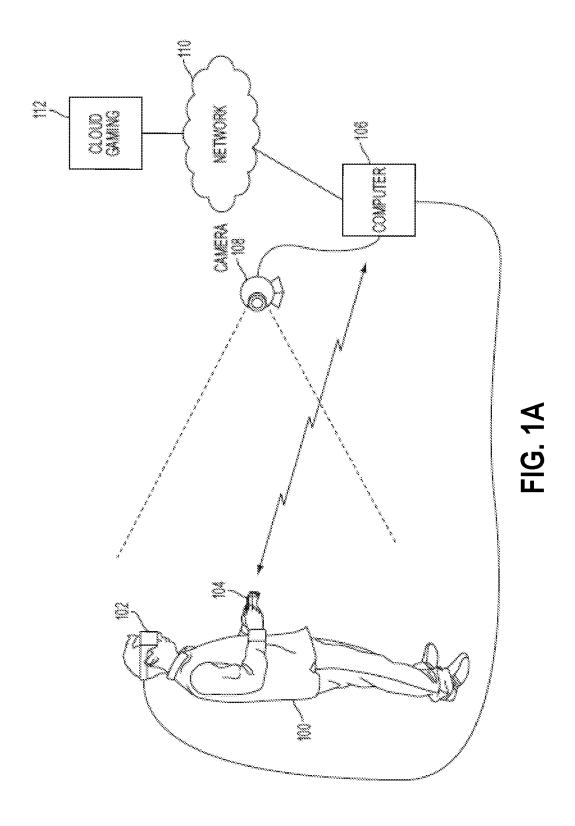
Lefohn, Aaron E. et al: "Resolution-Matched Shadow Maps", ACM Transactions on Graphics (TOG), ACM, US, vol. 26, No. 4, Oct. 1, 2007, 17 pages.

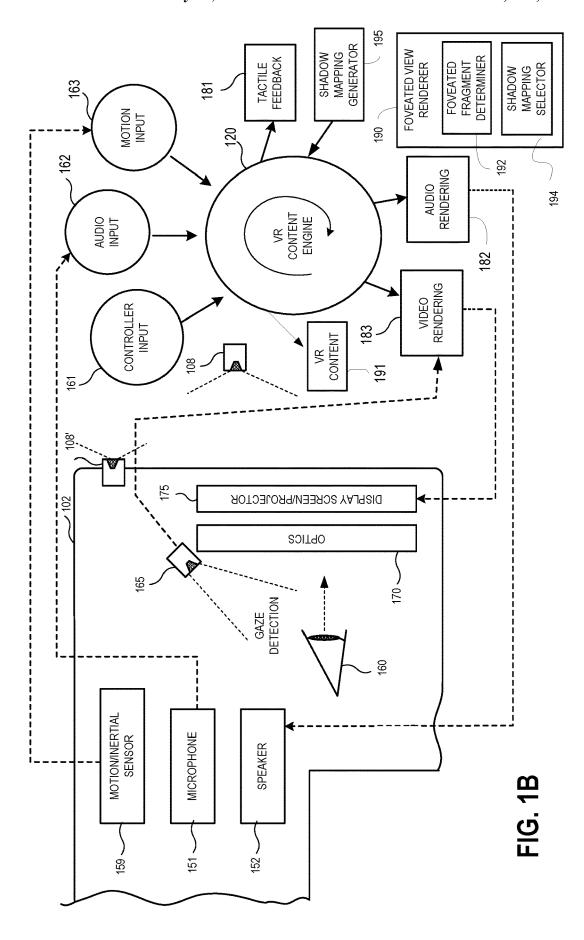
Huang, Peng et al: "Multi-resolution Shadow Mapping using CUDA rasterizer", 2013 International Conference on Computer-Aided Design and Computer Graphics, IEEE, Nov. 16, 2013, pp. 148-155.

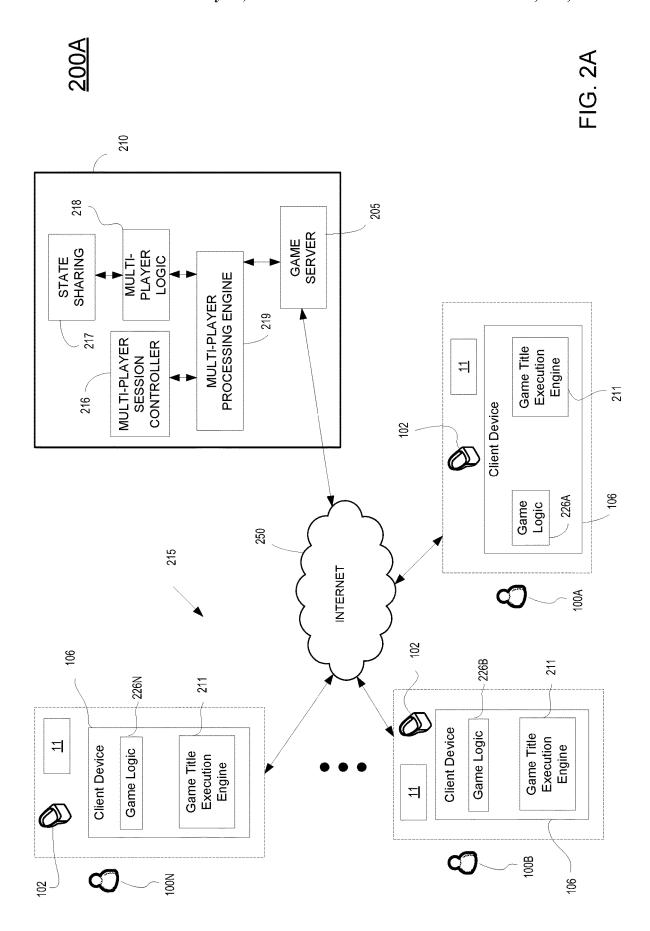
Cornudet, Veronique_Annex to Form PCT/ISA/206 Communication Relating to the Results of the Partial International Search, PCT/US2018/036675, dated Aug. 23, 2018, 11 pages.

Flohr, Chantal_Form PCT/ISA/220, ISA210, ISA237, PCT Notification of Transmittal of the International Search Report and the Written Opinion of the International Searching Authority, or the Declaration, PCT/US2018/036675, dated Nov. 20, 2018, 19 pages.

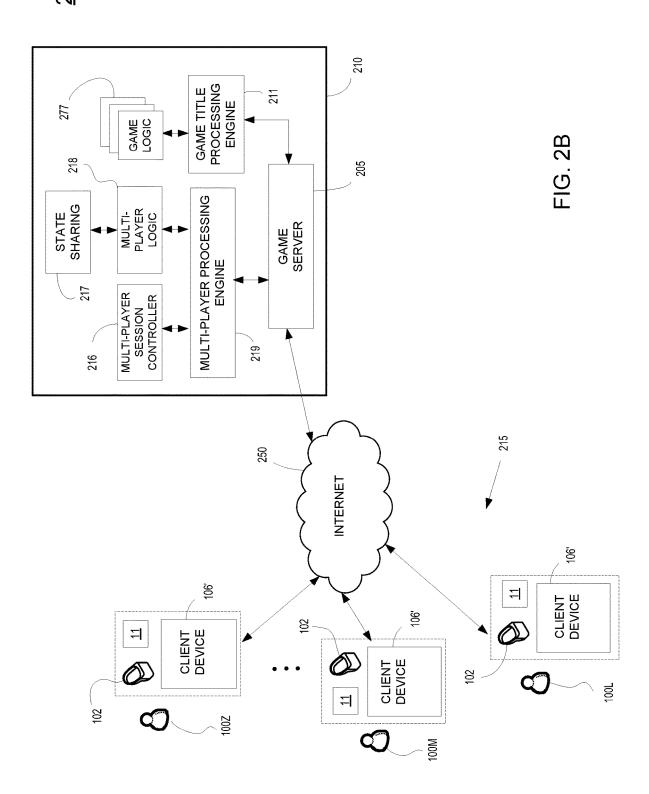
^{*} cited by examiner

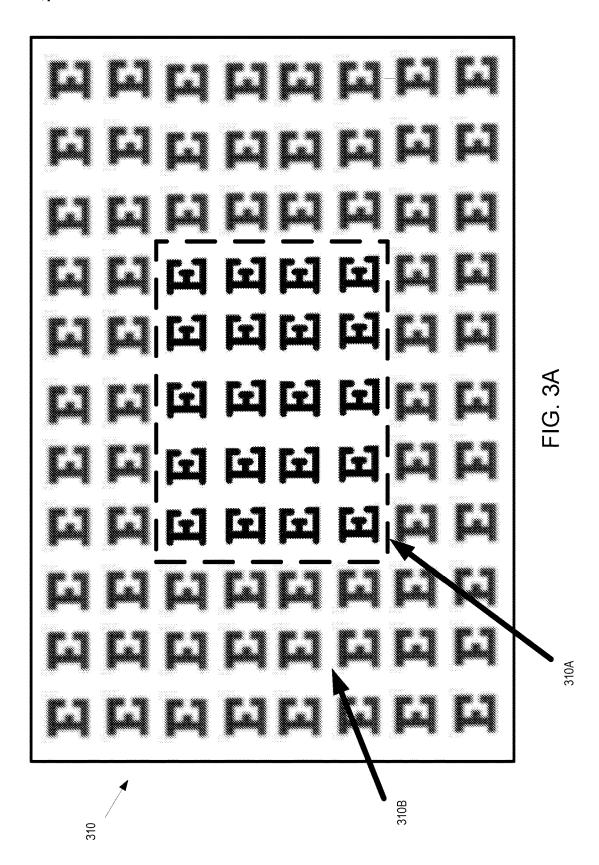


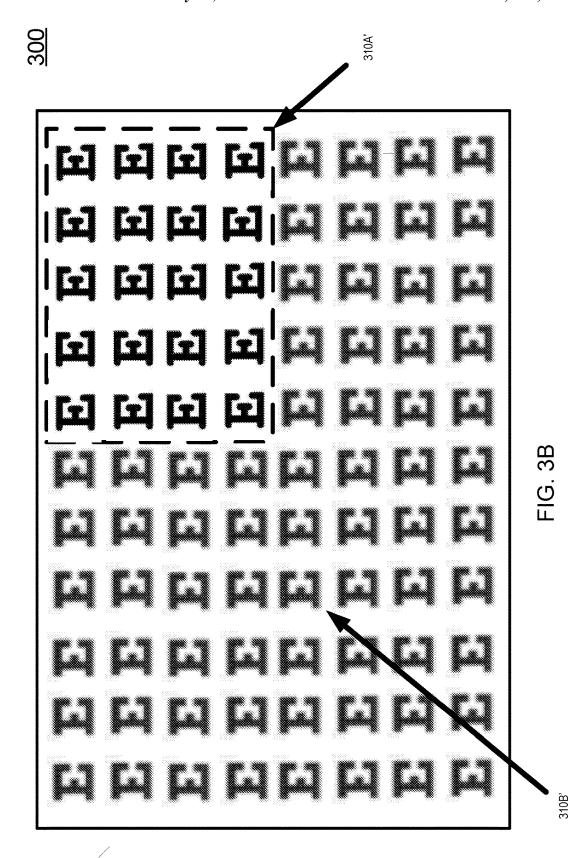




200B







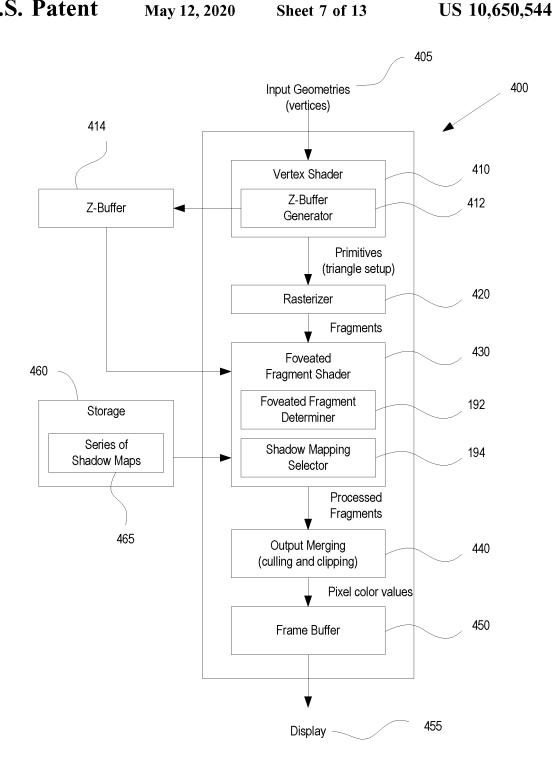


FIG. 4

<u>500</u>

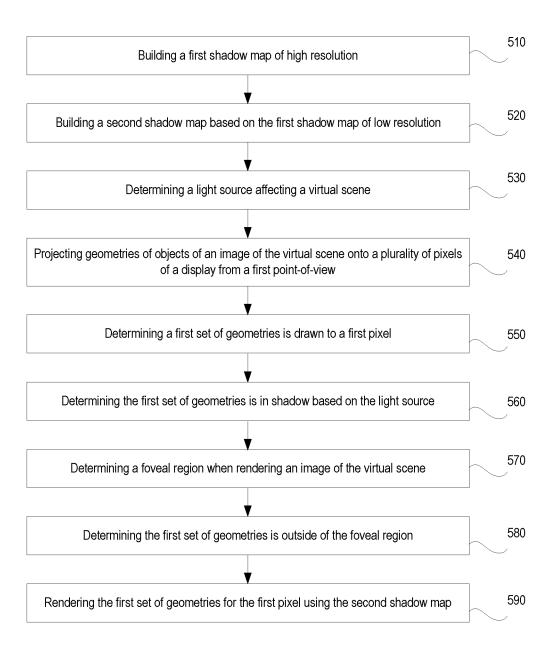
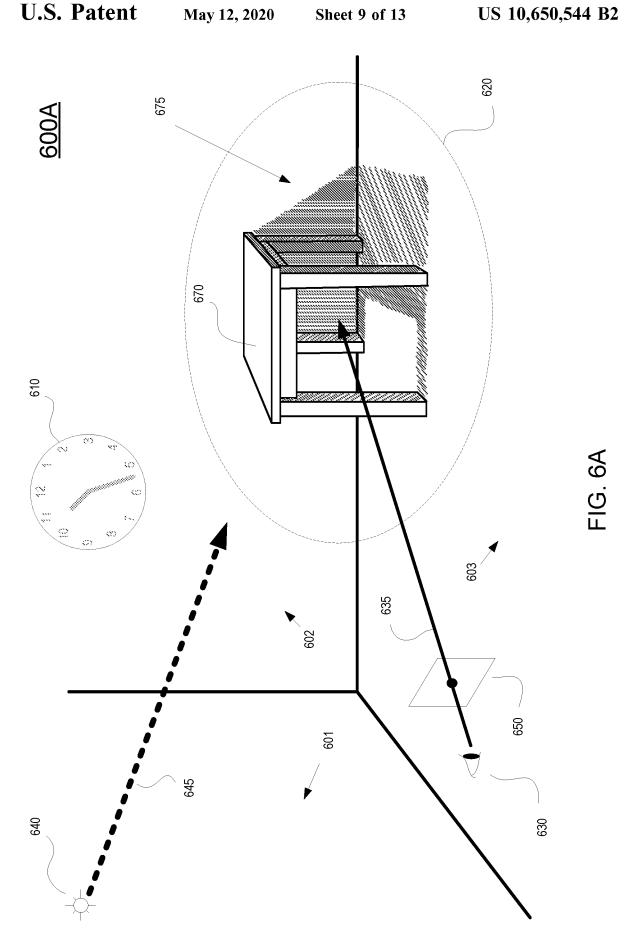
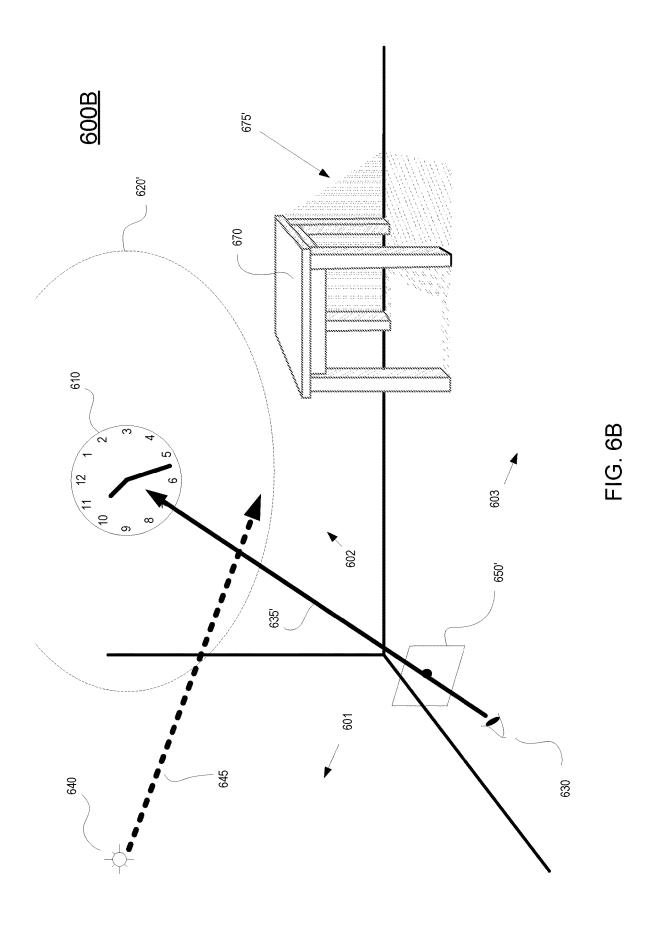
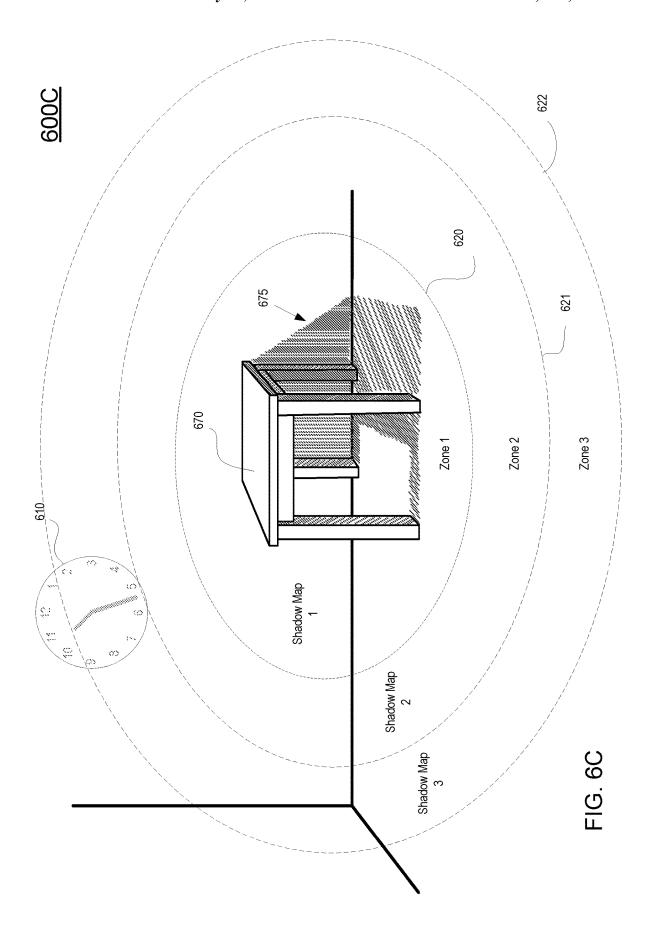


FIG. 5







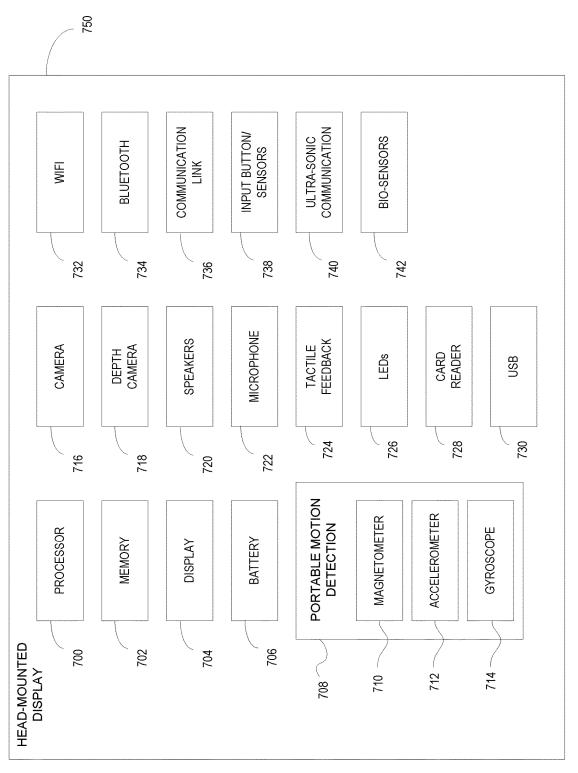


FIG. 7

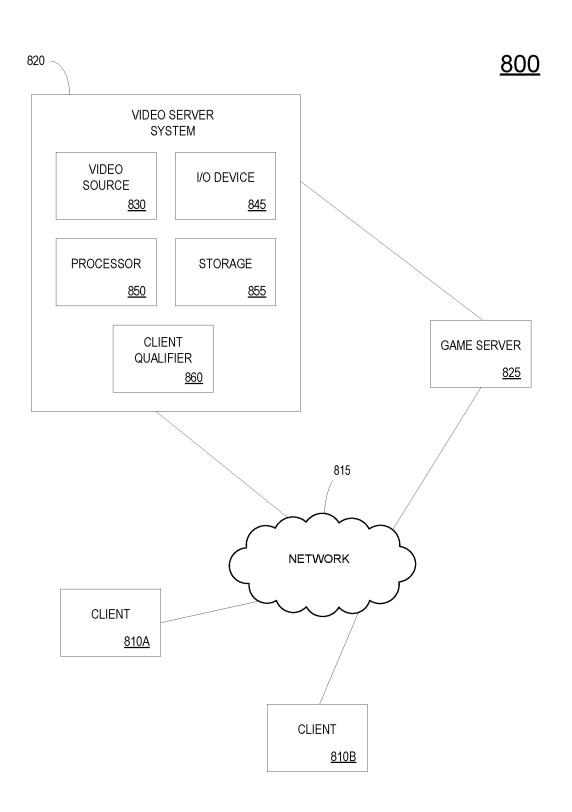


FIG. 8

OPTIMIZED SHADOWS IN A FOVEATED RENDERING SYSTEM

CLAIM OF PRIORITY

The present application claims priority to and the benefit of the commonly owned, provisional patent application, U.S. Ser. No. 62/517,829, filed on Jun. 9, 2017, entitled "OPTIMIZED SHADOWS IN A FOVEATED RENDERING SYSTEM," which is herein incorporated by reference in its entirety.

TECHNICAL FIELD

The present disclosure is related to video games or gaming applications. Among other things, this disclosure ¹⁵ describes methods and systems for selecting appropriate shadowing of objects in a foveated rendering system.

BACKGROUND OF THE DISCLOSURE

Video gaming has increasingly become more popular with the advancement of video game technology. For example, high powered graphics processors provide an unbelievably viewing and interactive experience when playing a video game. In addition, displays are being designed with higher and higher resolutions. For example, present technology includes displays having 2K resolution (e.g., 2.2 megapixels over 2048×1080 pixels) with an aspect ratio of approximately 19:10. Other displays having 4K UHD (Ultra high definition) resolution (e.g., 8.2 megapixels over 3840×2160 pixels) with an aspect ratio of 16:9 are now pushing into the market and is expected to gain traction. Increased graphics processing capabilities along with high resolution displays provide for a heretofore unbelievable viewing experience for the user, especially when playing a video game and gaming engine designed to take advantage of the higher resolution 35

Pushing rendered images/frames to a high resolution display also requires increased bandwidth capabilities, such as between the rendering engine and the display. In most cases, a wired connection should be able to handle the 40 required bandwidth supporting the display. However, gaming systems increasingly are configured with a wireless connection that may provide a bottleneck when pushing data to the display. For instance, a wireless connection may be established between a gaming console local to the user and the display. In these cases, the wireless connection may not be robust enough to handle the required bandwidth to fully take advantage of the higher resolution displays, such that the video game as displayed may be interrupted (as the buffer is filling up) in order to display the entire video sequence as rendered. In some cases, the video game pro- $^{50}\,$ cessing may be throttled in order to match the lower bandwidth of the wireless connection to the display, such that the video frames may be rendered at lower resolutions in order to push video data over the wireless connection without interruption; however, by throttling the processing, the user 55 is denied the full gaming experience with higher resolution

It would be beneficial to modify the graphics processing in order for the user to achieve a high level of satisfaction for the user, especially when playing a video game.

It is in this context that embodiments of the disclosure arise.

SUMMARY

Embodiments of the present disclosure relate to foveated rendering configured to display portions of images in a 2

foveal region with high resolution and portions outside the foveal region with lower resolution. In particular, shadowing of objects in an image that are inside the foveal region is rendered using a shadow map of high resolution. Shadowing of objects that are outside the foveal region is rendered using a shadow map of lower resolution. As such, instead of computing all shadows at a higher resolution, only shadows that are displayed within the foveal region are rendered using a shadow map of higher resolution. In that manner, the total bandwidth for the sequence of video frames being displayed is reduced, for example due in part to less complex images. Further, the sequence of video frames can be delivered (e.g., over wired or wireless connections) in real time with minimal or no latency because the computation complexity is reduced.

In one embodiment, a method for implementing a graphics pipeline is disclosed. The method includes building a first shadow map of high resolution. The method includes building a second shadow map based on the first shadow map, wherein the second shadow map has a lower resolution than the first shadow map. The method includes determining a light source affecting a virtual scene. The method includes projecting geometries of objects of an image of the virtual scene onto a plurality of pixels of a display from a first point-of-view. The method includes determining a foveal region when rendering the image of the virtual scene, wherein the foveal region corresponds to where an attention of a user is directed. The method includes determining a first set of geometries is drawn to a first pixel. The method includes determining the first set of geometries is in shadow based on the light source. The method includes determining the first set of geometries is outside of the foveal region. The method includes rendering the first set of geometries for the first pixel using the second shadow map.

In still another embodiment, a computer system is disclosed. The computer system including a processor and memory, wherein the memory is coupled to the processor and having stored therein instructions that, if executed by the computer system, cause the computer system to execute a method for implementing a graphics pipeline. The method includes building a first shadow map of high resolution. The method includes building a second shadow map based on the first shadow map, wherein the second shadow map has a lower resolution than the first shadow map. The method includes determining a light source affecting a virtual scene. The method includes projecting geometries of objects of an image of the virtual scene onto a plurality of pixels of a display from a first point-of-view. The method includes determining a foveal region when rendering the image of the virtual scene, wherein the foveal region corresponds to where an attention of a user is directed. The method includes determining a first set of geometries is drawn to a first pixel. The method includes determining the first set of geometries is in shadow based on the light source. The method includes determining the first set of geometries is outside of the foveal region. The method includes rendering the first set of geometries for the first pixel using the second shadow map.

In another embodiment, a non-transitory computer-readable medium storing a computer program for implementing a graphics pipeline is disclosed. The computer-readable medium includes program instructions for building a first shadow map of high resolution. The medium includes program instructions for building a second shadow map based on the first shadow map, wherein the second shadow map based on the first shadow map, wherein the second shadow map. The medium includes program instructions for determining a light source affecting a virtual scene. The medium includes

3

program instructions for projecting geometries of objects of an image of the virtual scene onto a plurality of pixels of a display from a first point-of-view. The medium includes program instructions for determining a foveal region when rendering the image of the virtual scene, wherein the foveal region corresponds to where an attention of a user is directed. The medium includes program instructions for determining a first set of geometries is drawn to a first pixel. The medium includes program instructions for determining the first set of geometries is in shadow based on the light source. The medium includes program instructions for determining the first set of geometries is outside of the foveal region. The medium includes program instructions for rendering the first set of geometries for the first pixel using the second shadow map.

Other aspects of the disclosure will become apparent from the following detailed description, taken in conjunction with the accompanying drawings, illustrating by way of example the principles of the disclosure.

BRIEF DESCRIPTION OF THE DRAWINGS

The disclosure may best be understood by reference to the following description taken in conjunction with the accompanying drawings in which:

FIG. 1A illustrates a system configured for providing an interactive experience with VR content, in accordance with one embodiment of the present disclosure.

FIG. 1B conceptually illustrates the function of a HMD in conjunction with an executing video game, in accordance 30 with an embodiment of the invention.

FIG. 2A illustrates a system providing gaming control to one or more users playing one or more gaming applications that are executing locally to the corresponding user, in accordance with one embodiment of the present disclosure. 35

FIG. 2B illustrates a system providing gaming control to one or more users playing a gaming application as executed over a cloud game network, in accordance with one embodiment of the present disclosure.

FIG. 3A illustrates an image shown on a display and 40 including a foveal region of high resolution, wherein the foveal region corresponds to a center of the display, in accordance with one embodiment of the present disclosure.

FIG. 3B illustrates an image shown on a display and including a foveal region of high resolution, wherein the 45 foveal region corresponds to a location of the display towards which the user is directing his or her gaze, in accordance with one embodiment of the present disclosure.

FIG. 4 illustrates a graphics processor implementing a graphics pipeline configured for foveated rendering including the rendering of shadows using different shadow maps, wherein a shadow is rendered using an appropriate shadow map based on whether the shadow is in the foveal region, and wherein the shadow is rendered using an appropriate shadow map based on its location as displayed in relation to 55 the foveal region, in accordance with one embodiment of the present disclosure.

FIG. 5 is a flow diagram illustrating steps in a method for implementing a graphics pipeline configured for foveated rendering including the rendering of shadows using different 60 shadow maps, wherein portions of an image in a foveal region are rendered with higher resolution, and portions of an image in a peripheral region are rendered with lower resolution, and wherein a shadow is rendered using an appropriate shadow map based on its location as displayed 65 in relation to the foveal region, in accordance with one embodiment of the disclosure.

4

FIG. 6A is an illustration of an image of a virtual scene of a gaming world, wherein the image shows the rendering of shadows as displayed in a foveal region, in accordance with one embodiment of the disclosure.

FIG. 6B is an illustration of the virtual scene introduced in FIG. 6A as rendered in an image and showing the rendering of shadows as displayed in a peripheral region, in accordance with one embodiment of the present disclosure.

FIG. 6C is an illustration of a series of shadow maps having progressively lower resolutions applied to the virtual scene introduced in FIG. 6A, in accordance with one embodiment of the present disclosure.

FIG. 7 is a diagram illustrating components of a headmounted display, in accordance with an embodiment of the 15 disclosure.

FIG. 8 is a block diagram of a Game System, according to various embodiments of the disclosure.

DETAILED DESCRIPTION

Although the following detailed description contains many specific details for the purposes of illustration, anyone of ordinary skill in the art will appreciate that many variations and alterations to the following details are within the scope of the present disclosure. Accordingly, the aspects of the present disclosure described below are set forth without any loss of generality to, and without imposing limitations upon, the claims that follow this description.

Generally speaking, the various embodiments of the present disclosure describe a graphics processor of a video rendering system that is configured to perform foveated rendering, wherein portions of images in a foveal region may be rendered with high resolution and portions outside the foveal region may be rendered with lower resolution. In particular, shadowing of objects that are displayed inside the foveal region are rendered using a shadow map of higher resolution, and the shadowing of objects that are displayed outside the foveal region are rendered using a shadow map of lower resolution. In some embodiments, the foveated rendering is performed within or for the purposes of displaying images within a head mounted display (HMD).

With the above general understanding of the various embodiments, example details of the embodiments will now be described with reference to the various drawings.

Throughout the specification, the reference to "video game" or "gaming application" is meant to represent any type of interactive application that is directed through execution of input commands. For illustration purposes only, an interactive application includes applications for gaming, word processing, video processing, video game processing, etc. Further, the terms video game and gaming application are interchangeable.

FIG. 1A illustrates a system for interactive gameplay of a video game, in accordance with an embodiment of the invention. A user 100 is shown wearing a head-mounted display (HMD) 102. The HMD 102 is worn in a manner similar to glasses, goggles, or a helmet, and is configured to display a video game from an interactive video game or other content from interactive application, to the user 100. The HMD 102 provides a very immersive experience to the user by virtue of its provision of display mechanisms in close proximity to the user's eyes. Thus, the HMD 102 can provide display regions to each of the user's eyes which occupy large portions or even the entirety of the field of view of the user. Though FIGS. 1A-1B and in other figures are shown using HMDs for displaying rendered images, embodiments of the present invention are well suited for

, in the second second

performing foveated rendering in any display device, wherein the foveated rendering includes rendering a shadow in an image based on the location of the shadow in relation to a foveal region, and displaying the rendered images on any display.

5

In one embodiment, HMD 102 is configurable to display images configured with foveated rendering, wherein portions of images in a foveal region are displayed with high resolution and portions outside the foveal region are displayed with lower resolution. In particular, shadowing is rendered differently depending on whether the shadow is inside or outside the foveal region as displayed. For example, shadowing of objects in an image that are inside the foveal region is rendered using a shadow map of high resolution. Shadowing of objects that are outside the foveal 15 region is rendered using a shadow map of lower resolution. As such, instead of computing all shadowing using a shadow map of higher resolution, only shadows that are displayed within the foveal region are rendered using a shadow map of higher resolution, thereby lessening the computation 20 required for rendering all shadowing in a corresponding image. Also, each image containing shadows rendered at lower resolution outside the foveal region require less defining data than images containing shadows all rendered at a higher resolution. In that manner, the total bandwidth for the 25 sequence of video frames being displayed is reduced.

In one embodiment, the HMD 102 can be connected to a computer or gaming console 106. The connection to computer 106 can be wired or wireless. The computer 106 can be any general or special purpose computer known in the art, 30 including but not limited to, a gaming console, personal computer, laptop, tablet computer, mobile device, cellular phone, tablet, thin client, set-top box, media streaming device, etc. In one embodiment, the computer 106 can be configured to execute a video game, and output the video 35 and audio from the video game for rendering by the HMD **102**. The computer **106** is not restricted to executing a video game but may also be configured to execute an interactive application, which outputs VR content 191 for rendering by the HMD 102. In one embodiment, computer 106 performs 40 the selection and use of shadow maps of lower resolution when rendering shadows of objects in an image that is displayed outside of the foveal region.

The user 100 may operate a controller 104 to provide input for the video game. Additionally, a camera 108 can be 45 configured to capture one or more images of the interactive environment in which the user 100 is located. These captured images can be analyzed to determine the location and movements of the user 100, the HMD 102, and the controller **104**. In one embodiment, the controller **104** includes a light 50 or other marker elements which can be tracked to determine its location and orientation. The camera 108 can include one or more microphones to capture sound from the interactive environment. Sound captured by a microphone array may be processed to identify the location of a sound source. Sound 55 from an identified location can be selectively utilized or processed to the exclusion of other sounds not from the identified location. Furthermore, the camera 108 can be defined to include multiple image capture devices (e.g. stereoscopic pair of cameras), an IR camera, a depth camera, 60 and combinations thereof.

In another embodiment, the computer 106 functions as a thin client in communication over a network with a cloud gaming provider 112. The cloud gaming provider 112 maintains and executes the video game being played by the user 65 102. The computer 106 transmits inputs from the HMD 102, the controller 104 and the camera 108, to the cloud gaming

6

provider, which processes the inputs to affect the game state of the executing video game. The output from the executing video game, such as video data, audio data, and haptic feedback data, is transmitted to the computer 106. The computer 106 may further process the data before transmission or may directly transmit the data to the relevant devices. For example, video and audio streams are provided to the HMD 102, whereas the haptic feedback data is used to generate a vibration feedback command, which is provided to the controller 104.

In one embodiment, the HMD 102, controller 104, and camera 108, may themselves be networked devices that connect to the network 110 to communicate with the cloud gaming provider 112. For example, the computer 106 may be a local network device, such as a router, that does not otherwise perform video game processing, but facilitates passage network traffic. The connections to the network by the HMD 102, controller 104, and camera (i.e., image capture device) 108 may be wired or wireless.

In vet another embodiment, the computer 106 may execute a portion of the video game, while the remaining portion of the video game may be executed on a cloud gaming provider 112. In other embodiments, portions of the video game may also be executed on HMD 102. For example, a request for downloading the video game from the computer 106 may be serviced by the cloud gaming provider 112. While the request is being serviced, the cloud gaming provider 112 may execute a portion of the video game and provide game content to the computer 106 for rendering on the HMD 102. The computer 106 may communicate with the cloud gaming provider 112 over a network 110. Inputs received from the HMD 102, the controller 104 and the camera 108, are transmitted to the cloud gaming provider 112, while the video game is downloading on to the computer 106. The cloud gaming provider 112 processes the inputs to affect the game state of the executing video game. The output from the executing video game, such as video data, audio data, and haptic feedback data, is transmitted to the computer 106 for onward transmission to the respective

Once the video game has been completely downloaded to the computer 106, the computer 106 may execute the video game and resume game play of the video game from where it was left off on the cloud gaming provider 112. The inputs from the HMD 102, the controller 104, and the camera 108 are processed by the computer 106, and the game state of the video game is adjusted, in response to the inputs received from the HMD 102, the controller 104, and the camera 108. In such embodiments, a game state of the video game at the computer 106 is synchronized with the game state at the cloud gaming provider 112. The synchronization may be done periodically to keep the state of the video game current at both the computer 106 and the cloud gaming provider 112. The computer 106 may directly transmit the output data to the relevant devices. For example, video and audio streams are provided to the HMD 102, whereas the haptic feedback data is used to generate a vibration feedback command, which is provided to the controller 104.

FIG. 1B conceptually illustrates the function of a HMD 102 in conjunction with the generation of VR content 191 (e.g., execution of an application and/or video game, etc.), in accordance with an embodiment of the invention. In some implementations, the VR content engine 120 is being executed on a computer 106 (not shown) that is communicatively coupled to the HMD 102. The computer may be local to the HMD (e.g., part of local area network) or may be remotely located (e.g., part of a wide area network, a

cloud network, etc.) and accessed via a network. The communication between the HMD 102 and the computer 106 may follow a wired or a wireless connection protocol. For example, the VR content engine 120 executing an application may be a video game engine executing a video game, 5 and is configured to receive inputs to update a game state of the video game. The following description of FIG. 1B is described within the context of the VR content engine 120 executing a video game, for purposes of brevity and clarity, and is intended to represent the execution of any application 10 capable of generating VR content 191. The game state of the video game can be defined, at least in part, by values of various parameters of the video game which define various aspects of the current gameplay, such as the presence and location of objects, the conditions of a virtual environment, 15 the triggering of events, user profiles, view perspectives, etc.

In the illustrated embodiment, the game engine 120 receives, by way of example, controller input 161, audio input 162 and motion input 163. The controller input 161 may be defined from the operation of a gaming controller 20 separate from the HMD 102, such as a hand-held gaming controller 104 (e.g. Sony DUALSHOCK®4 wireless controller, Sony PlayStation® Move motion controller) or wearable controllers, such as wearable glove interface controller, etc. By way of example, controller input 161 may include 25 directional inputs, button presses, trigger activation, movements, gestures or other kinds of inputs processed from the operation of a gaming controller. The audio input 162 can be processed from a microphone 151 of the HMD 102, or from a microphone included in the image capture device 108 or 30 elsewhere within the local system environment. The motion input 163 can be processed from a motion sensor 159 included in the HMD 102, or from image capture device 108 as it captures images of the HMD 102. The VR content engine 120 (e.g., executing a gaming application) receives 35 inputs which are processed according to the configuration of the game engine to update the game state of the video game. The engine 120 outputs game state data to various rendering modules which process the game state data to define content which will be presented to the user.

In the illustrated embodiment, a video rendering module 183 is defined to render a video stream for presentation on the HMD 102. Foveated view renderer 190 is configured to render foveated images in conjunction with and/or independent of video rendering module 183. Additionally, the func- 45 tionality provided by the foveated view renderer 190 may be incorporated within the video rendering module 183, in embodiments. In particular, foveated view renderer 190 is configured to perform foveated rendering, wherein portions of images in a foveal region are rendered with high resolution and portions outside the foveal region are rendered with lower resolution. More particularly, shadowing of objects that are outside the foveal region is rendered using a shadow map of lower resolution in order to reduce the computation for rendering a corresponding image. Shadowing of objects 55 that are inside the foveal region is rendered using a shadow map of higher resolution. The foveated fragment determiner 192 is configured to determine whether a fragment being rendered is displayed inside or outside a foveal region. The shadow map selector 194 is configured to determine which 60 shadow map to use when rendering a corresponding fragment. For example, when the fragment is displayed inside the foveal region, a shadow map of higher resolution is selected, and when the fragment is displayed outside the foveal region, a shadow map of lower resolution is selected. 65 Further, the shadow map selector 194 is configured to select an appropriate shadow map based on the distance the

8

corresponding fragment is from the foveal region, as displayed, wherein a series of shadow maps are generated based on distances from the foveal region such that resolutions of shadow maps in the series decrease as the distance from the foveal region increases. Shadow mapping generator 195 is configured to build the shadow maps of higher and lower resolution, and including the series of shadow maps described above. For example, a first shadow map may be generated at higher resolution, and other mipmaps of the first shadow map may be generated based on the first shadow map (e.g., shadow maps of decreasing resolutions). In particular, a series of mipmaps may define shadow maps each having progressively lower resolutions of the first shadow map.

A lens of optics 170 in the HMD 102 is configured for viewing the VR content 191. A display screen 175 is disposed behind the lens of optics 170, such that the lens of optics 170 is between the display screen 175 and an eye 160 of the user, when the HMD 102 is worn by the user. In that manner, the video stream may be presented by the display screen/projector mechanism 175, and viewed through optics 170 by the eye 160 of the user. An HMD user may elect to interact with the interactive VR content 191 (e.g., VR video source, video game content, etc.) by wearing the HMD and selecting a video game for game play, for example. Interactive virtual reality (VR) scenes from the video game are rendered on the display screen 175 of the HMD. In that manner, the HMD allows the user to completely immerse in the game play by provisioning display mechanism of the HMD in close proximity to the user's eyes. The display regions defined in the display screen of the HMD for rendering content may occupy large portions or even the entirety of the field of view of the user. Typically, each eye is supported by an associated lens of optics 170 which is viewing one or more display screens.

An audio rendering module 182 is configured to render an audio stream for listening by the user. In one embodiment, the audio stream is output through a speaker 152 associated with the HMD 102. It should be appreciated that speaker 152 may take the form of an open air speaker, headphones, or any other kind of speaker capable of presenting audio.

In one embodiment, a gaze tracking camera 165 is included in the HMD 102 to enable tracking of the gaze of the user. Although only one gaze tracking camera 165 is included, it should be noted that more than one gaze tracking camera may be employed to track the gaze of the user. The gaze tracking camera captures images of the user's eyes, which are analyzed to determine the gaze direction of the user. In one embodiment, information about the gaze direction of the user can be utilized to affect the video rendering. For example, if a user's eyes are determined to be looking in a specific direction, then the video rendering for that direction can be prioritized or emphasized, such as by providing greater detail, higher resolution through foveated rendering as provided by foveated view renderer 190, higher resolution of shadowing displayed in the foveal region, lower resolution of shadowing displayed outside the foveal region, or faster updates in the region where the user is looking. It should be appreciated that the gaze direction of the user can be defined relative to the head mounted display, relative to a real environment in which the user is situated, and/or relative to a virtual environment that is being rendered on the head mounted display.

Broadly speaking, analysis of images captured by the gaze tracking camera 165, when considered alone, provides for a gaze direction of the user relative to the HMD 102. However, when considered in combination with the tracked

location and orientation of the HMD 102, a real-world gaze direction of the user can be determined, as the location and orientation of the HMD 102 is synonymous with the location and orientation of the user's head. That is, the real-world gaze direction of the user can be determined from tracking 5 the positional movements of the user's eyes and tracking the location and orientation of the HMD 102. When a view of a virtual environment is rendered on the HMD 102, the real-world gaze direction of the user can be applied to determine a virtual world gaze direction of the user in the 10 virtual environment.

Additionally, a tactile feedback module 181 is configured to provide signals to tactile feedback hardware included in either the HMD 102 or another device operated by the HMD user, such as a controller 104. The tactile feedback may take 15 the form of various kinds of tactile sensations, such as vibration feedback, temperature feedback, pressure feedback, etc.

FIG. 2A illustrates a system 200A providing gaming control to one or more users playing one or more gaming 20 applications that are executing locally to the corresponding user, and wherein back-end server support (e.g., accessible through game server 205) may be configured to support a plurality of local computing devices supporting a plurality of users, wherein each local computing device may be execut- 25 ing an instance of a video game, such as in a single-player or multi-player video game. For example, in a multi-player mode, while the video game is executing locally, the cloud game network concurrently receives information (e.g., game state data) from each local computing device and distributes 30 that information accordingly throughout one or more of the local computing devices so that each user is able to interact with other users (e.g., through corresponding characters in the video game) in the gaming environment of the multiplayer video game. In that manner, the cloud game network 35 coordinates and combines the game plays for each of the users within the multi-player gaming environment. Referring now to the drawings, like referenced numerals designate identical or corresponding parts.

As shown in FIG. 2A, a plurality of users 215 (e.g., user 40 100A, user 100B . . . user 100N) is playing a plurality of gaming applications, wherein each of the gaming applications is executed locally on a corresponding client device 106 (e.g., game console) of a corresponding user. Each of the client devices 106 may be configured similarly in that local 45 execution of a corresponding gaming application is performed. For example, user 100A may be playing a first gaming application on a corresponding client device 106, wherein an instance of the first gaming application is executed by a corresponding game title execution engine 50 211. Game logic 226A (e.g., executable code) implementing the first gaming application is stored on the corresponding client device 106, and is used to execute the first gaming application. For purposes of illustration, game logic may be delivered to the corresponding client device 106 through a 55 cates with multi-player logic 218 in order to enable interportable medium (e.g., flash drive, compact disk, etc.) or through a network (e.g., downloaded through the internet 250 from a gaming provider). In addition, user 100B is playing a second gaming application on a corresponding client device 106, wherein an instance of the second gaming 60 application is executed by a corresponding game title execution engine 211. The second gaming application may be identical to the first gaming application executing for user 100A or a different gaming application. Game logic 226B (e.g., executable code) implementing the second gaming 65 application is stored on the corresponding client device 106 as previously described, and is used to execute the second

10

gaming application. Further, user 100N is playing an Nth gaming application on a corresponding client device 106, wherein an instance of the Nth gaming application is executed by a corresponding game title execution engine 211. The Nth gaming application may be identical to the first or second gaming application, or may be a completely different gaming application. Game logic 226N (e.g., executable code) implementing the third gaming application is stored on the corresponding client device 106 as previously described, and is used to execute the Nth gaming application.

As previously described, client device 106 may receive input from various types of input devices, such as game controllers, tablet computers, keyboards, gestures captured by video cameras, mice touch pads, etc. Client device 106 can be any type of computing device having at least a memory and a processor module that is capable of connecting to the game server 205 over network 150. Also, client device 106 of a corresponding user is configured for generating rendered images executed by the game title execution engine 211 executing locally or remotely, and for displaying the rendered images on a display (e.g., display 11, HMD 102, etc.). For example, the rendered images may be associated with an instance of the first gaming application executing on client device 106 of user 100A. For example, a corresponding client device 106 is configured to interact with an instance of a corresponding gaming application as executed locally or remotely to implement a game play of a corresponding user, such as through input commands that are used to drive game play. Some examples of client device 106 include a personal computer (PC), a game console, a home theater device, a general purpose computer, mobile computing device, a tablet, a phone, or any other types of computing devices that can interact with the game server 205 to execute an instance of a video game.

In one embodiment, client device 106 is operating in a single-player mode for a corresponding user that is playing a gaming application. In another embodiment, multiple client devices 106 are operating in a multi-player mode for corresponding users that are each playing a specific gaming application. In that case, back-end server support via the game server may provide multi-player functionality, such as through the multi-player processing engine 219. In particular, multi-player processing engine 219 is configured for controlling a multi-player gaming session for a particular gaming application. For example, multi-player processing engine 219 communicates with the multi-player session controller 216, which is configured to establish and maintain communication sessions with each of the users and/or players participating in the multi-player gaming session. In that manner, users in the session can communicate with each other as controlled by the multi-player session controller

Further, multi-player processing engine 219 communiaction between users within corresponding gaming environments of each user. In particular, state sharing module 217 is configured to manage states for each of the users in the multi-player gaming session. For example, state data may include game state data that defines the state of the game play (of a gaming application) for a corresponding user at a particular point. For example, game state data may include game characters, game objects, game object attributes, game attributes, game object state, graphic overlays, etc. In that manner, game state data allows for the generation of the gaming environment that exists at the corresponding point in the gaming application. Game state data may also include

the state of every device used for rendering the game play, such as states of CPU, GPU, memory, register values, program counter value, programmable DMA state, buffered data for the DMA, audio chip state, CD-ROM state, etc. Game state data may also identify which parts of the 5 executable code need to be loaded to execute the video game from that point. Game state data may be stored in database (not shown), and is accessible by state sharing module 217.

Further, state data may include user saved data that includes information that personalizes the video game for 10 the corresponding player. This includes information associated with the character played by the user, so that the video game is rendered with a character that may be unique to that user (e.g., location, shape, look, clothing, weaponry, etc.). In that manner, the user saved data enables generation of a 15 character for the game play of a corresponding user, wherein the character has a state that corresponds to the point in the gaming application experienced currently by a corresponding user. For example, user saved data may include the game difficulty selected by a corresponding user 100A when 20 playing the game, game level, character attributes, character location, number of lives left, the total possible number of lives available, armor, trophy, time counter values, etc. User saved data may also include user profile data that identifies a corresponding user 100A, for example. User saved data 25 may be stored in storage (not shown).

In that manner, the multi-player processing engine 219 using the state sharing data 217 and multi-player logic 218 is able to overlay/insert objects and characters into each of the gaming environments of the users participating in the 30 multi-player gaming session. For example, a character of a first user is overlaid/inserted into the gaming environment of a second user. This allows for interaction between users in the multi-player gaming session via each of their respective gaming environments (e.g., as displayed on a screen).

FIG. 2B illustrates a system 200B providing gaming control to one or more users 215 (e.g., users 100L, 100M...100Z) playing a gaming application in respective VR viewing environments as executed over a cloud game network, in accordance with one embodiment of the present 40 disclosure. In some embodiments, the cloud game network may be a game cloud system 210 that includes a plurality of virtual machines (VMs) running on a hypervisor of a host machine, with one or more virtual machines configured to execute a game processor module utilizing the hardware 45 resources available to the hypervisor of the host. Referring now to the drawings, like referenced numerals designate identical or corresponding parts.

As shown, the game cloud system 210 includes a game server 205 that provides access to a plurality of interactive 50 video games or gaming applications. Game server 205 may be any type of server computing device available in the cloud, and may be configured as one or more virtual machines executing on one or more hosts. For example, game server 205 may manage a virtual machine supporting 55 a game processor that instantiates an instance of a gaming application for a user. As such, a plurality of game processors of game server 205 associated with a plurality of virtual machines is configured to execute multiple instances of the gaming application associated with game plays of the plu- 60 rality of users 215. In that manner, back-end server support provides streaming of media (e.g., video, audio, etc.) of game plays of a plurality of gaming applications to a plurality of corresponding users.

A plurality of users 215 accesses the game cloud system 65 210 via network 250, wherein users (e.g., users 100L, 100M . . . 100Z) access network 250 via corresponding

client devices 106', wherein client device 106' may be configured similarly as client device 106 of FIG. 2A (e.g., including game executing engine 211, etc.), or may be configured as a thin client providing that interfaces with a back end server providing computational functionality (e.g., including game executing engine 211).

In particular, a client device 106' of a corresponding user 100L is configured for requesting access to gaming applications over a network 250, such as the internet, and for rendering instances of gaming application (e.g., video game) executed by the game server 205 and delivered to a display device associated with the corresponding user 100L. For example, user 100L may be interacting through client device 106' with an instance of a gaming application executing on game processor of game server 205. More particularly, an instance of the gaming application is executed by the game title execution engine 211. Game logic (e.g., executable code) implementing the gaming application is stored and accessible through a data store (not shown), and is used to execute the gaming application. Game title processing engine 211 is able to support a plurality of gaming applications using a plurality of game logics 277, as shown.

As previously described, client device 106' may receive input from various types of input devices, such as game controllers, tablet computers, keyboards, gestures captured by video cameras, mice touch pads, etc. Client device 106' can be any type of computing device having at least a memory and a processor module that is capable of connecting to the game server 205 over network 250. Also, client device 106' of a corresponding user is configured for generating rendered images executed by the game title execution engine 211 executing locally or remotely, and for displaying the rendered images on a display. For example, the rendered images may be associated with an instance of 35 the first gaming application executing on client device 106' of user 100L. For example, a corresponding client device 106' is configured to interact with an instance of a corresponding gaming application as executed locally or remotely to implement a game play of a corresponding user, such as through input commands that are used to drive game play.

Client device 106' is configured for receiving rendered images, and for displaying the rendered images on display 11 and/or HMD 102 (e.g., displaying VR content). For example, through cloud based services the rendered images may be delivered by an instance of a gaming application executing on game executing engine 211 of game server 205 in association with user 100. In another example, through local game processing, the rendered images may be delivered by the local game executing engine 211. In either case, client device 106 is configured to interact with the local or remote executing engine 211 in association with the game play of a corresponding user 100, such as through input commands that are used to drive game play. In another implementation, the rendered images may be streamed to a smartphone or tablet, wirelessly or wired, direct from the cloud based services or via the client device 106 (e.g., PlayStation® Remote Play).

In another embodiment, multi-player processing engine 219, previously described, provides for controlling a multi-player gaming session for a gaming application. In particular, when the multi-player processing engine 219 is managing the multi-player gaming session, the multi-player session controller 216 is configured to establish and maintain communication sessions with each of the users and/or players in the multi-player session. In that manner, users in the session can communicate with each other as controlled by the multi-player session controller 216.

Further, multi-player processing engine 219 communicates with multi-player logic 218 in order to enable interaction between users within corresponding gaming environments of each user. In particular, state sharing module 217 is configured to manage states for each of the users in the 5 multi-player gaming session. For example, state data may include game state data that defines the state of the game play (of a gaming application) for a corresponding user 100 at a particular point, as previously described. Further, state data may include user saved data that includes information 10 that personalizes the video game for the corresponding player, as previously described. For example, state data includes information associated with the user's character, so that the video game is rendered with a character that may be unique to that user (e.g., shape, look, clothing, weaponry, 15 etc.). In that manner, the multi-player processing engine 219 using the state sharing data 217 and multi-player logic 218 is able to overlay/insert objects and characters into each of the gaming environments of the users participating in the multi-player gaming session. This allows for interaction 20 between users in the multi-player gaming session via each of their respective gaming environments (e.g., as displayed on a screen).

FIG. 3A illustrates an image 310 shown on a display 300, wherein the image includes a foveal region 310A of high 25 resolution, wherein the foveal region corresponds to a center of the display, in accordance with one embodiment of the present disclosure. In particular, image 310 includes rows and columns of the letter "E" for simplicity and clarity. The image 310 is partitioned into multiple regions, including a 30 foveal region 310A and a peripheral region 310B.

As shown, the foveal region 310A is static and corresponds to the center of the display 300. The foveal region 310A is assumed to be the region towards which the user mostly directs his or her gaze (e.g., using the fovea of the 35 eye), such as when viewing graphics of a video game. Though the gaze of the user may occasionally be directed off center, the gaze is mostly directed to the center (to view the main content). In some cases, the image is designed to initially bring the gaze of the user off-center (e.g., to view an 40 object of interest), but then to bring the gaze back to the center (e.g., by moving the object towards the foveal region 310A).

In particular, the portion or portions of any image, such as image 310, as displayed and located within the foveal region 45 310A will be rendered at higher resolution. For example, the graphics pipeline will render portions of the image in the foveal region 310A while minimizing the use of any techniques used to reduce computational complexity. In particular, for embodiments of the present invention, light sources 50 affecting objects displayed using pixels corresponding to the foveal region 310A are individually computed within the graphics pipeline in order to determine each of their effects on the objects (e.g., color, texture, shadowing, etc. on polygons of the objects). Representative of the higher reso- 55 lution, the letter "E" objects as displayed within the foveal region 310A are shown with clarity, vibrant color, and minimal blurriness. This is consistent with and takes advantage of the gaze of the user being directed towards the foveal region 310A on display 300.

In addition, the portion or portions of an image, such as image 310, as disposed and located in the peripheral region 310B will be rendered at lower resolution (e.g., lower than the resolution of the portions of the image and/or objects located in the foveal region 310A). The gaze of the user is 65 not typically directed to the objects located in and/or displayed in the peripheral region 310B, as the main focus of

14

the gaze is directed to the objects in the foveal region 310A. Consistent with real-life views into a scene, the objects in the peripheral region 310B are rendered with lower resolution, with enough detail so that the user is able to perceive moving objects (e.g., a human walking straight-legged instead of with knees bent) and with sufficient contrast within an object or between objects in the peripheral region 310B, for example. To achieve rendering at lower resolutions, the graphics pipeline may render portions of the image in the peripheral region 310B using computationally efficient techniques that reduce computational complexity. In particular, for embodiments of the present invention, shadowing of objects in an image that are inside the foveal region is rendered using a shadow map of high resolution. Shadowing of objects that are outside the foveal region is rendered using a shadow map of lower resolution. As such, instead of computing all shadows at a higher resolution, only shadows that are displayed within the foveal region are rendered using a shadow map of higher resolution. This reduces computational processing when rendering the overall image, and especially for objects rendered in the peripheral region 310B.

FIG. 3B illustrates an image 310' shown on a display 300 and including a foveal region 310A' of high resolution, wherein the foveal region corresponds to a location of the display towards which the user is directing his or her gaze, in accordance with one embodiment of the present disclosure. In particular, image 310' is similar to image 310 shown in FIG. 3A and includes rows and columns of the letter "E" for simplicity and clarity. The image 310 is partitioned into multiple regions, including a foveal region 310A' and a peripheral region 310B'.

As shown, the foveal region 310A' is dynamically moving throughout display 300 depending on which direction the gaze of the user is directed towards. As previously described, the gaze may be tracked using gaze tracking camera 165 of HMD 102, for example. As such, the foveal region 310A' may not necessarily correspond to the center of display 300, but instead correlates to the actual direction and focus of attention of the user within image 310'. That is, the foveal region 310A' dynamically moves with the movement of the eye and/or eyes of the user.

As previously introduced, the portion or portions of any image, such as image 310', as displayed and located within the foveal region 310A' will be rendered at higher resolution by minimizing the use of any rendering techniques used to reduce computational complexity when rendering objects located in the foveal region 310A'. In particular, for embodiments of the present invention, light sources affecting objects displayed using pixels corresponding to the foveal region 310A' are individually computed within the graphics pipeline in order to determine each of their effects on the objects (e.g., color, texture, shadowing, etc. on polygons of the objects). Representative of the higher resolution, the letter "E" objects as displayed within the foveal region 310A' are shown with clarity, vibrant color, and minimal blurriness.

In addition, the portion or portions of an image, such as image 310', as disposed and located in the peripheral region 310B' will be rendered at lower resolution (e.g., lower than the resolution of the portions of the image and/or objects located in the foveal region 310A). As previously introduced, the gaze of the user is not typically directed to the objects located in and/or displayed in the peripheral region 310B', as the main focus of the gaze is directed to the objects in the foveal region 310A'. As such, the objects in the peripheral region 310B' are rendered with lower resolution,

with enough detail so that the user is able to perceive moving objects (e.g., a human walking straight-legged instead of with knees bent) and with sufficient contrast within an object or between objects in the peripheral region 310B', for example. To achieve rendering at lower resolutions, the 5 graphics pipeline may render portions of the image in the peripheral region 310B' using computationally efficient techniques that reduce computational complexity. In particular, for embodiments of the present invention, shadowing of objects that are in the peripheral region 310B' is 10 rendered using a shadow map of lower resolution.

FIG. 4 illustrates a graphics processor implementing a graphics pipeline 400A configured for foveated rendering, in accordance with one embodiment of the present disclosure. The graphics pipeline 400A is illustrative of the general 15 process for rendering images using 3D (three dimensional) polygon rendering processes, but includes an additional programmable element within the pipeline that performs foveated rendering. The graphics pipeline 400A for a rendered image outputs corresponding color information for 20 each of the pixels in a display, wherein the color information may represent texture and shading (e.g., color, shadowing, etc.). Graphics pipeline 400A is implementable within the game console 106 of FIG. 1A, VR content engine 120 of FIG. 1B, client devices 106 of FIGS. 2A and 2B, and/or 25 game title processing engine 211 of FIG. 2B.

As shown, the graphics pipeline receives input geometries 405. For example, the input geometries 405 may include vertices within a 3D gaming world, and information corresponding to each of the vertices. A given object within the 30 gaming world can be represented using polygons (e.g., triangles) defined by vertices, wherein the surface of a corresponding polygon is then processed through the graphics pipeline 400A to achieve a final effect (e.g., color, texture, etc.). Vertex attributes may include normal (e.g., 35 which direction is the light in relation to the vertex), color (e.g., RGB—red, green, and blue triple, etc.), and texture coordinate/mapping information.

The vertex shader and/or program 410 receives the input geometries 405, and builds the polygons or primitives that 40 make up the objects within the 3D scene. That is, the vertex shader 410 builds up the objects using the primitives as they are placed within the gaming world. The vertex shader 410 may be configured to perform lighting and shadowing calculations for the polygons, which is dependent on the 45 lighting for the scene. The primitives are output by the vertex shader 410 and delivered to the next stage of the graphics pipeline 400A. Additional operations may also be performed by the vertex shader 410 such as clipping (e.g., identify and disregard primitives that are outside the viewing 50 frustum as defined by the viewing location in the gaming world).

In embodiments of the present invention, the graphics processor 400 is configured to perform foveated rendering, including rendering shadows in an efficient manner Shadow 55 mapping is implemented within graphics processor 400 wherein shadows of objects in the virtual scene are rendered in an image of the scene using shadow maps. As previously described, a series of shadow maps are created based on a first shadow map generated at the highest resolution. The 60 series of shadow maps, including the first shadow map, includes a pre-determined sequence of shadow images, each of which is associated with a progressively lower resolution (e.g., based on the first shadow map). The shadow mapping generator 195 of FIG. 1B is configured to generate the first shadow map. In addition, the shadow mapping generator 195 is configured to generate the related shadow maps of

lower resolution using mipmap techniques. Generally, the first shadow map of highest resolution is duplicated at a reduced level of detail to generate a second shadow map of lower resolution. Depending on the reduction applied, multiple shadow maps of decreasing resolution may be generated, each of which is based on the first shadow map. Using shadow maps of reduced resolution increases the rendering speed of an image, since the number of texels used to render the shadow in the image is smaller than when rendering shadows at higher resolutions.

The shadow maps may be stored in storage 460. In particular, a series of shadow maps 465 may be located in storage 460 after generation by the shadow mapping generator 195, and accessed by the graphics pipeline 400 in order to properly render the shadow using a shadow map having an appropriate resolution. In one embodiment, a binary system of shadow maps is implemented as the series of shadow maps 465, wherein the first shadow map is of higher resolution, and a second shadow map is of a resolution that is lower than the first shadow map. More particularly, the first shadow map of higher resolution is used for rendering shadows located in a foveal region of an image as displayed, and the second shadow map of lower resolution is used for rendering shadows located in a peripheral region. In another embodiment, the series of shadow maps is built based on distances from a foveal region, wherein resolutions of shadow maps in the series decrease as the distance from the foveal region increases.

The vertex shader 410 includes a z-buffer generator 412 which is configured to generate a depth or z-buffer 414 of values. In particular, the virtual scene is rendered from the viewpoint of a corresponding light source in order to determine which objects and/or surfaces cast a shadow and which objects and/or surfaces are in shadow. As the scene is rendered depth texture is captured in the z-buffer 414, which stores the depth of the closest surfaces in the virtual scene for each pixel in a reference image plane. For illustration, along a line from the light source that passes through a corresponding pixel, multiple objects may be rendered along the line at various depths. The z-buffer 414 stores the depth that is closest to the light source, which corresponds to the surface that is casting a shadow along that line. The z-buffer 414 is referenced when rendering polygons by the foveated fragment shader and/or program 430.

The primitives output by the vertex shader 410 are fed into the rasterizer 420 that is configured to project objects in the scene to a two-dimensional (2D) image plane defined by the viewing location in the 3D gaming world (e.g., camera location, user eye location, etc.). At a simplistic level, the rasterizer 420 looks at each primitive and determines which pixels are affected by the corresponding primitive. In particular, the rasterizer 420 partitions the primitives into pixel sized fragments, wherein each fragment corresponds to a pixel in the display and/or a reference plane associated with the rendering point-of-view (e.g., camera view). It is important to note that one or more fragments may contribute to the color of a corresponding pixel when displaying an image. Additional operations may also be performed by the rasterizer 420 such as clipping (identify and disregard fragments that are outside the viewing frustum) and culling (disregard fragments that are occluded by closer objects) to the viewing location.

The foveated fragment shader 430 at its core performs shading operations on the fragments to determine how the color and brightness of a primitive varies with available lighting. For example, fragment shader 430 may determine depth, color, normal and texture coordinates (e.g., texture

details) for each fragment, and may further determine appropriate levels of light, darkness, and color for the fragments. In particular, fragment shader 430 calculates the traits of each fragment, including color and other attributes (e.g., z-depth for distance from the viewing location, and alpha values for transparency). In addition, the fragment shader 430 applies lighting effects to the fragments based on the available lighting affecting the corresponding fragments. Further, the fragment shader 430 may apply shadowing effects for each fragment, as described below. For purposes of description only light sources are described as being point lights having a definite position in the gaming world, and which radiates light omni-directionally. Other lighting is available, such as directional lighting, etc.

More particularly, the foveated fragment shader 430 performs shading operations as described above based on whether the fragment is within the foveal region or peripheral region. Fragments that are located within the foveal region of the displayed image are processed using shading operations at high resolution, without regard to processing efficiency in order to achieve detailed texture and color values for fragments within the foveal region. On the other hand, the foveated fragment shader 430 performs shading operations on fragments that are located within the peripheral region with an interest in processing efficiency in order to process fragments with sufficient detail with minimal operations, such as providing movement and sufficient contrast.

For example, in embodiments of the present invention 30 shadowing of objects that are displayed inside a foveal region are rendered using a shadow map of higher resolution, and the shadowing of objects that are displayed in the peripheral region are rendered using a shadow map of lower resolution. In particular, the foveated fragment shader is 35 configured to determine whether a fragment is in shadow, and/or determine whether a plurality of fragments of one or more objects contributing to the color of a pixel is in shadow by accessing the values in the z-buffer 414. Correlation of the coordinates of the fragment in 3D gaming space and 40 view space (e.g., from camera point-of-view, or viewing location) may be performed by the vertex shader 410 or fragment shader 430. In particular, once it is determined that a fragment is in shadow, the shadow mapping selector 194 selects the appropriate shadow map depending on the loca- 45 tion of the fragment. The foveated fragment determiner 192 is configured to determine whether the fragment is displayed in the foveal region or in the peripheral region. For instance, when the fragment is displayed in the foveal region, the shadow mapping selector 194 chooses the first shadow map 50 of highest resolution from the series of shadow maps 465 to render the fragment. Also, when the fragment is displayed in the non-foveated reign, the shadow mapping selector 194 chooses the second shadow map of lower resolution from the series of shadow maps 465 to render the fragment. This 55 may be an illustration of binary shadow mapping the series 465. In addition, the shadow mapping selector 194 may choose a shadow map from the series of shadow maps 465 having an appropriate resolution that is based on the distance from the foveal region that the fragment is rendered and 60 displayed. As the distance increases and goes further away from the foveal region, shadow maps of lower resolution are

The output of the fragment shader 430 includes processed fragments (e.g., texture and shading information to include shadowing) and is delivered to the next stage of the graphics pipeline $400\mathrm{A}$.

18

The output merging component 440 calculates the traits of each pixel depending on the fragments that contribute and/or affect each corresponding pixel. That is, the fragments of all primitives in the 3D gaming world are combined into the 2D color pixel for the display. For example, fragments that contribute to texture and shading information for a corresponding pixel are combined to output a final color value for the pixel delivered to the next stage in the graphics pipeline 400A. The output merging component 440 may perform optional blending of values between fragments and/or pixels determined from the fragment shader 430.

Color values for each pixel in the display are stored in the frame buffer 455. These values are scanned to the corresponding pixels when displaying a corresponding image of the scene. In particular, the display reads color values from the frame buffer for each pixel, row-by-row, from left-to-right or right-to-left, top-to-bottom or bottom-to-top, or any other pattern, and illuminates pixels using those pixel values when displaying the image.

With the detailed description of the various modules of the gaming server and client device communicating over a network, flow diagram 500 of FIG. 5 discloses a method for implementing a graphics pipeline configured to perform foveated rendering, wherein shadows displayed inside the foveal region are rendered using a shadow map of higher resolution, and the shadows displayed in the non-foveated are rendered using a shadow map of lower resolution, in accordance with one embodiment of the present disclosure. Flow diagram 500 is implemented within a client device or cloud based gaming system when rendering images, as previously described. In one implementation, the client device is a gaming console (also referred to as game console).

At **510**, the method includes building and/or generating a first shadow map of high resolution. This may be the shadow map of highest resolution, which is used for generating shadow maps of lower resolution (e.g., mipmapping). That is, at **520**, the method includes building and/or generating a second shadow map based on the first shadow map, wherein the second shadow map has a lower resolution than the first shadow map. This may be an illustration of a binary system of shadow maps. Further, a series of shadow maps may be generated based on the first shadow map, wherein the series includes shadow maps of progressively lower resolution.

At 530, the method includes determining a light source affecting a virtual scene. The light source may be one of multiple light sources, each of which affect a virtual scene. The lighting from the light sources will determine the final color of corresponding pixels when displaying a corresponding object within the scene as viewed from a viewing location. When determining shadowing, each of the light sources must be considered to fully render the shadowing contained within an image. For illustration, a single light source is used to disclose the embodiments of the present invention. In one embodiment, the same process used to determine shadowing is performed for each of the multiple light sources and combined and/or merged to form the final coloring of pixels in the display.

At **540**, the method includes projecting geometries of objects of an image of the scene onto a plurality of pixels of a display from a first point-of-view (e.g., the viewing location). As previously introduced, the rasterizer and/or vertex shader are configured to associate primitives and/or fragments of primitives with corresponding pixels. Specifically, the vertex shader may be configured to associate primitives with corresponding pixels, wherein before rasterization the scene by primitives is partitioned into tiles. In

addition, the rasterizer may be configured to associate fragments of primitives with corresponding pixels, wherein the scene by fragments may be partitioned into tiles.

At **550**, the method includes determining a first set of geometries is drawn to a first pixel. The set of geometries 5 may include one or more fragments of one or more objects. For example, the vertex shader may be configured to determine which pixels are affected by or drawn by the fragments in the set, as previously described.

At **560**, the method includes determining that the first set 10 of geometries is in shadow based on the light source. As previously described, a depth test is performed by the fragment shader to determine which fragments are in shadow from the light source. The depth test is applied to each of the set of geometries associated with the pixel. For 15 example, when in shadow, geometries in the set will have a depth that is greater than the depth stored in the z-buffer (which is the depth of the closest object to the light source for fragments contributing to that pixel). As such, each of the set of geometries is in shadow.

At 570, the method includes determining a foveal region when rendering the image of the virtual scene, wherein the foveal region corresponds to where an attention of a user is directed. The foveal region may be associated with rendering surfaces at the highest resolution. The foveal region may 25 be where the attention of the user is focused, either by assumption (e.g., generally at the center of the display to cover a static area of the display) or by eye tracking (e.g., to determine the direction of the gaze). For a static foveal region, the foveal region is centered to the center of the 30 display. When performing eye tracking, the gaze of the user viewing the virtual scene is tracked, wherein the gaze indicates a direction into the virtual scene towards which the attention of the user is directed. In this case, the foveal region may be centered in the image based on the direction 35 of the gaze. In addition, the portion of the image (e.g., primitives and/or fragments) in the foveal region is rendered by the graphics processor at high resolution. For example, shadowing in the foveal region is rendered using shadow maps of higher and/or the highest resolution. On the other 40 hand, the peripheral region does not have the attention of the user and the portion of the image (e.g., primitives and/or fragments) in the peripheral region is rendered by the graphics processor at a lower resolution. For example, shadowing in the peripheral region is rendered using shadow 45 maps of lower resolution. More particularly, a first subset of pixels is assigned to the foveal region, and wherein primitives and/or fragments associated with the first subset is processed at high resolution. A second subset of pixels is assigned to a peripheral region that is outside of the foveal 50 region, and wherein primitives and/or fragments associated with the second subset is processed at the lower resolution (e.g., lower than the resolution associated with the foveal region).

In addition, at **580**, the method includes determining that 55 the first set of geometries is outside of the foveal region. That is, the first set of geometries affects a pixel of a display that is in the peripheral region. In this case, for efficient rendering, at **590**, the method includes rendering the first set of geometries for the first pixel using the second shadow 60 map of lower resolution. In that manner, fewer computations are required to render the shadowing as projected by the first set of geometries.

Further, for geometries that are inside the foveal region, shadowing is rendered using a shadow map of higher 65 resolution. For example, it may be determined that a second set of geometries is drawn to a second pixel. Further, the

second set of geometries is determined to be inside the foveal region towards which the attention of the user is directed. As such, the second set of geometries for the second pixel may be rendered for the second pixel using the

20

second pixel may be rendered for the second first shadow map of higher resolution.

In addition, other shadow effects may be disabled when rendering the first set of geometries for the first pixel. That is, to reduce the computational cost, shadowing in the peripheral region is limited to rendering using the second shadow map of lower resolution. Beyond the use of the second shadow map, no further shadow effects, or minimal shadow effects, are computed within the graphics pipeline.

In still another embodiment, a series of shadow maps is generated, wherein the shadow maps is based on the shadow map of highest resolution, and resolutions of the shadow maps in the series progressively decrease. The series of shadow maps is applied when rendering fragments based on the distance between a corresponding fragment as displayed and the foveal region. For example, a third set of geometries 20 may be determined to be drawn to a third pixel. The third pixel is determined to be outside the foveal region. In addition, the third set of geometries may be determined to be in shadow from the light source based on z-buffer testing, as previously described. A distance is determined between the third set of geometries as drawn in the third pixel and the foveal region. A third shadow map is selected from the series of shadow maps based on the determined distance. Basically, the further the pixel is when rendering the third set of geometries from the foveal region, the lower the resolution of the shadow map selected. After selection, the third set of geometries is rendered for the third pixel using the third shadow map, which has a lower resolution than the first shadow map of highest resolution, and may have a lower resolution than a second shadow map.

FIG. 6A is an illustration of an image 600A of a virtual scene of a gaming world, wherein the image shows the rendering of shadows as displayed in a foveal region, in accordance with one embodiment of the disclosure. The image 600A shows a room having walls 601, 602, and a floor 603. Two objects are shown in the room, to include an object 670 (a table) resting on floor 603, and a clock 610 hanging on wall 602.

For purposes of brevity and clarity, a light source **640** is shown in the image **600**A that affects the virtual scene. As shown, light source **640** has a location in the gaming world (e.g., the room). The location may be defined by a coordinate system (not shown). The light source **640** may be directional or omnidirectional, which in either case will cast a shadow **675**. As shown, light source **640** casts a shadow **675** of table **670** along at least one direction **645** (wherein the direction **645** is superimposed onto image **600**A).

The image 600A may be rendered from a particular viewing location within the gaming world, and displayed. For example, the image may be rendered from a virtual image plane associated with a viewing location in the gaming world. For instance, an eye 630 is superimposed onto image 600A to illustrate a gaze of a user along a direction 635. Image 600A may be projected within image plane 650 that is centered on direction 635. As shown in FIG. 6A, the gaze of the user is directed towards table 670 along direction 635 in image 600A.

In addition, the attention of the user is focused along direction 635. As such, a foveal region 620 of the image as displayed is defined based on the direction of the gaze of the user (e.g., assumed and/or tracked). As previously described, objects in the foveal region 620 are rendered at high resolution. In addition, in embodiments of the present invention,

shadowing inside the foveal region **620** is rendered using a shadow map of higher resolution (e.g., the first shadow map of highest resolution from which other shadow maps are generated—such as using mipmaps). For example, shadow **675** created by light source **640** shining on table **670** is 5 rendered using a high resolution shadow map.

As shown, objects outside of the foveal region are rendered at lower resolution, as previously described. For example, clock 610 is displayed outside of the foveal region 620, and is rendered at a lower resolution. For illustration, 10 clock 610 is faintly shown to represent rendering at lower resolution. That is, because clock 610 is in a peripheral region, less detail is required to render clock 610 in image 600A.

FIG. **6**B is an illustration of the virtual scene introduced 15 in FIG. **6**A as rendered in an image **600**B and showing the rendering of shadows as displayed in a peripheral region, in accordance with one embodiment of the present disclosure

The image 600B is similar to the image 600A except the attention of the user is towards a different object. As 20 previously introduced, image 600B shows a room having walls 601, 602, and a floor 603. Two objects are shown in the room, to include an object 670 (a table) resting on floor 603, and a clock 610 hanging on wall 602. The light source 640 is shown in image 600A, and casts a shadow 675' of table 25 670 along at least one direction 645 (wherein the direction 645 is superimposed onto image 600B).

In addition, the attention of the user is now focused along direction 635'. As such, a foveal region 620' of the image 600B as displayed is defined based on the direction of the 30 gaze of the user (e.g., assumed and/or tracked). As previously described, objects in the foveal region 620' are rendered at high resolution. As such, in FIG. 6B, the clock 610 is now rendered at high resolution, and is clearly and boldly shown in FIG. 6B.

As shown, objects outside of the foveal region 620' are rendered at lower resolution, as previously described. In FIG. $6\mathrm{B}$, table 670 is now faintly shown to represent rendering at lower resolution. Because the table 670 is in a peripheral region, less detail is required to render table 670 40 in image $600\mathrm{B}$.

In addition, in embodiments of the present invention, shadowing outside the foveal region 620' is rendered using a shadow map of lower resolution (e.g., the second or succeeding shadow map of lower resolution). For example, 45 shadow 675' created by light source 640 shining on table 670 is rendered using a lower resolution shadow map. Shadow 675' is shown with less detail than shadow 675 of FIG. 6A, wherein shadow 675' does not require full detail as it is located in a peripheral region.

FIG. 6C is an illustration of a series of shadow maps having progressively lower resolutions applied to the virtual scene introduced in FIG. 6A, in accordance with one embodiment of the present disclosure. As previously described, the series of shadow maps is generated based on 55 the shadow map of highest resolution, and resolutions of the shadow maps in the series progressively decrease. In one embodiment, the series of shadow maps is applied when rendering fragments based on the distance between a corresponding fragment as displayed and the foveal region. As 60 shown in FIG. 6C, the foveal region 620 includes objects that are rendered at high resolution. Foveal region 620 is also referred to as Zone 1 in FIG. 6C. Because the table 670 is located in Zone 1, the foveal region, table 670 is rendered in great detail—at high resolution. In addition, shadow 675 is rendered using shadow map 1 of the highest resolution in the series of shadow maps because shadow 675 is located in

22

Zone 1. Objects shown in the foveal region 620 have a minimum or zero distance from the foveal region 620, and as such, the shadow map of highest resolution is selected for rendering shadows.

In addition, dotted boundary 621 and the outer boundary of foveal region 620 define Zone 2, which is the nearest zone that is outside the foveal region 620. Shadows in Zone 2 are rendered using shadow map 2, which is of a lower resolution than that of shadow map 1. Further, dotted boundary 621 and dotted boundary 622 define Zone 3, which is also outside the foveal region 620. Zone 3 is further from the foveal region 620 than Zone 2. As such, shadows in Zone 3 are rendered using shadow map 3, which is a lower resolution than that of shadow map 2. In the series shadow map 3 has the lowest resolution and is associated with Zone 3, which is furthest from the foveal region.

While specific embodiments have been provided to demonstrate the implementation of a graphics processor of a video rendering system that is configured to perform fove-ated rendering, wherein portions of images in a foveal region are rendered with high resolution (e.g., rendering shadows using a shadow map of higher resolution), and portions in the peripheral region are rendered with lower resolution (e.g., rendering shadows using a shadow map of lower resolution), these are described by way of example and not by way of limitation. Those skilled in the art having read the present disclosure will realize additional embodiments falling within the spirit and scope of the present disclosure.

It should be noted, that access services, such as providing access to games of the current embodiments, delivered over a wide geographical area often use cloud computing. Cloud computing is a style of computing in which dynamically scalable and often virtualized resources are provided as a service over the Internet. Users do not need to be an expert in the technology infrastructure in the "cloud" that supports them. Cloud computing can be divided into different services, such as Infrastructure as a Service (IaaS), Platform as a Service (PaaS), and Software as a Service (SaaS). Cloud computing services often provide common applications, such as video games, online that are accessed from a web browser, while the software and data are stored on the servers in the cloud. The term cloud is used as a metaphor for the Internet, based on how the Internet is depicted in computer network diagrams and is an abstraction for the complex infrastructure it conceals.

A Game Processing Server (GPS) (or simply a "game server") is used by game clients to play single and multiplayer video games. Most video games played over the Internet operate via a connection to the game server. Typically, games use a dedicated server application that collects data from players and distributes it to other players. This is more efficient and effective than a peer-to-peer arrangement, but it requires a separate server to host the server application. In another embodiment, the GPS establishes communication between the players and their respective game-playing devices to exchange information without relying on the centralized GPS.

Dedicated GPSs are servers which run independently of the client. Such servers are usually run on dedicated hardware located in data centers, providing more bandwidth and dedicated processing power. Dedicated servers are the preferred method of hosting game servers for most PC-based multiplayer games. Massively multiplayer online games run on dedicated servers usually hosted by a software company that owns the game title, allowing them to control and update content.

Users access the remote services with client devices, which include at least a CPU, a display and I/O. The client device can be a PC, a mobile phone, a netbook, a PDA, etc. In one embodiment, the network executing on the game server recognizes the type of device used by the client and 5 adjusts the communication method employed. In other cases, client devices use a standard communications method, such as html, to access the application on the game server over the internet.

Embodiments of the present disclosure may be practiced 10 with various computer system configurations including hand-held devices, microprocessor-based or programmable consumer electronics, minicomputers, mainframe computers and the like. The disclosure can also be practiced in distributed computing environments 15 where tasks are performed by remote processing devices that are linked through a wire-based or wireless network.

It should be appreciated that a given video game may be developed for a specific platform and a specific associated controller device. However, when such a game is made 20 available via a game cloud system as presented herein, the user may be accessing the video game with a different controller device. For example, a game might have been developed for a game console and its associated controller, whereas the user might be accessing a cloud-based version 25 of the game from a personal computer utilizing a keyboard and mouse. In such a scenario, the input parameter configuration can define a mapping from inputs which can be generated by the user's available controller device (in this case, a keyboard and mouse) to inputs which are acceptable 30 for the execution of the video game.

In another example, a user may access the cloud gaming system via a tablet computing device, a touchscreen smartphone, or other touchscreen driven device. In this case, the client device and the controller device are integrated 35 together in the same device, with inputs being provided by way of detected touchscreen inputs/gestures. For such a device, the input parameter configuration may define particular touchscreen inputs corresponding to game inputs for the video game. For example, buttons, a directional pad, or 40 other types of input elements might be displayed or overlaid during running of the video game to indicate locations on the touchscreen that the user can touch to generate a game input. Gestures such as swipes in particular directions or specific touch motions may also be detected as game inputs. In one 45 embodiment, a tutorial can be provided to the user indicating how to provide input via the touchscreen for gameplay, e.g. prior to beginning gameplay of the video game, so as to acclimate the user to the operation of the controls on the touchscreen.

In some embodiments, the client device serves as the connection point for a controller device. That is, the controller device communicates via a wireless or wired connection with the client device to transmit inputs from the controller device to the client device. The client device may 55 in turn process these inputs and then transmit input data to the cloud game server via a network (e.g., accessed via a local networking device such as a router). However, in other embodiments, the controller can itself be a networked device, with the ability to communicate inputs directly via 60 the network to the cloud game server, without being required to communicate such inputs through the client device first. For example, the controller might connect to a local networking device (such as the aforementioned router) to send to and receive data from the cloud game server. Thus, while 65 the client device may still be required to receive video output from the cloud-based video game and render it on a local

display, input latency can be reduced by allowing the controller to send inputs directly over the network to the cloud game server, bypassing the client device.

In one embodiment, a networked controller and client device can be configured to send certain types of inputs directly from the controller to the cloud game server, and other types of inputs via the client device. For example, inputs whose detection does not depend on any additional hardware or processing apart from the controller itself can be sent directly from the controller to the cloud game server via the network, bypassing the client device. Such inputs may include button inputs, joystick inputs, embedded motion detection inputs (e.g., accelerometer, magnetometer, gyroscope), etc. However, inputs that utilize additional hardware or require processing by the client device can be sent by the client device to the cloud game server. These might include captured video or audio from the game environment that may be processed by the client device before sending to the cloud game server. Additionally, inputs from motion detection hardware of the controller might be processed by the client device in conjunction with captured video to detect the position and motion of the controller, which would subsequently be communicated by the client device to the cloud game server. It should be appreciated that the controller device in accordance with various embodiments may also receive data (e.g., feedback data) from the client device or directly from the cloud gaming server.

It should be understood that the embodiments described herein may be executed on any type of client device. In some embodiments, the client device is a head mounted display (HMD).

FIG. 7, a diagram illustrating components of a headmounted display 102 is shown, in accordance with an embodiment of the disclosure. The head-mounted display 102 includes a processor 700 for executing program instructions. A memory 702 is provided for storage purposes, and may include both volatile and non-volatile memory. A display 704 is included which provides a visual interface that a user may view. A battery 706 is provided as a power source for the head-mounted display 102. A motion detection module 708 may include any of various kinds of motion sensitive hardware, such as a magnetometer 710A, an accelerometer 712, and a gyroscope 714.

An accelerometer is a device for measuring acceleration and gravity induced reaction forces. Single and multiple axis models are available to detect magnitude and direction of the acceleration in different directions. The accelerometer is used to sense inclination, vibration, and shock. In one embodiment, three accelerometers 712 are used to provide the direction of gravity, which gives an absolute reference for two angles (world-space pitch and world-space roll).

A magnetometer measures the strength and direction of the magnetic field in the vicinity of the head-mounted display. In one embodiment, three magnetometers 710A are used within the head-mounted display, ensuring an absolute reference for the world-space yaw angle. In one embodiment, the magnetometer is designed to span the earth magnetic field, which is ±80 microtesla. Magnetometers are affected by metal, and provide a yaw measurement that is monotonic with actual yaw. The magnetic field may be warped due to metal in the environment, which causes a warp in the yaw measurement. If necessary, this warp can be calibrated using information from other sensors such as the gyroscope or the camera. In one embodiment, accelerometer 712 is used together with magnetometer 710A to obtain the inclination and azimuth of the head-mounted display 102.

A gyroscope is a device for measuring or maintaining orientation, based on the principles of angular momentum. In one embodiment, three gyroscopes 714 provide information about movement across the respective axis (x, y and z) based on inertial sensing. The gyroscopes help in detecting 5 fast rotations. However, the gyroscopes can drift overtime without the existence of an absolute reference. This requires resetting the gyroscopes periodically, which can be done using other available information, such as positional/orientation determination based on visual tracking of an object, 10 accelerometer, magnetometer, etc.

A camera 716 is provided for capturing images and image streams of a real environment. More than one camera may be included in the head-mounted display 102, including a camera that is rear-facing (directed away from a user when 15 the user is viewing the display of the head-mounted display 102), and a camera that is front-facing (directed towards the user when the user is viewing the display of the head-mounted display 102). Additionally, a depth camera 718 may be included in the head-mounted display 102 for 20 sensing depth information of objects in a real environment.

In one embodiment, a camera integrated on a front face of the HMD may be used to provide warnings regarding safety. For example, if the user is approaching a wall or object, the user may be warned. In one embodiment, the use may be 25 provided with an outline view of physical objects in the room, to warn the user of their presence. The outline may, for example, be an overlay in the virtual environment. In some embodiments, the HMD user may be provided with a view to a reference marker, that is overlaid in, for example, 30 the floor. For instance, the marker may provide the user a reference of where the center of the room is, which in which the user is playing the game. This may provide, for example, visual information to the user of where the user should move to avoid hitting a wall or other object in the room. Tactile 35 warnings can also be provided to the user, and/or audio warnings, to provide more safety for when the user wears and plays games or navigates content with an HMD.

The head-mounted display 102 includes speakers 720 for providing audio output. Also, a microphone 722 may be 40 included for capturing audio from the real environment, including sounds from the ambient environment, speech made by the user, etc. The head-mounted display 102 includes tactile feedback module 724 for providing tactile feedback to the user. In one embodiment, the tactile feedback module 724 is capable of causing movement and/or vibration of the head-mounted display 102 so as to provide tactile feedback to the user.

LEDs **726** are provided as visual indicators of statuses of the head-mounted display **102**. For example, an LED may 50 indicate battery level, power on, etc. A card reader **728** is provided to enable the head-mounted display **102** to read and write information to and from a memory card. A USB interface **730** is included as one example of an interface for enabling connection of peripheral devices, or connection to 55 other devices, such as other portable devices, computers, etc. In various embodiments of the head-mounted display **102**, any of various kinds of interfaces may be included to enable greater connectivity of the head-mounted display **102**.

A Wi-Fi module 732 is included for enabling connection 60 to the Internet via wireless networking technologies. Also, the head-mounted display 102 includes a Bluetooth module 734 for enabling wireless connection to other devices. A communications link 736 may also be included for connection to other devices. In one embodiment, the communications link 736 utilizes infrared transmission for wireless communication. In other embodiments, the communications

26

link 736 may utilize any of various wireless or wired transmission protocols for communication with other devices

Input buttons/sensors 738 are included to provide an input interface for the user. Any of various kinds of input interfaces may be included, such as buttons, touchpad, joystick, trackball, etc. An ultra-sonic communication module 740 may be included in head-mounted display 102 for facilitating communication with other devices via ultra-sonic technologies.

Bio-sensors 742 are included to enable detection of physiological data from a user. In one embodiment, the biosensors 742 include one or more dry electrodes for detecting bio-electric signals of the user through the user's skin.

Photo-sensors **744** are included to respond to signals from emitters (e.g., infrared base stations) placed in a 3-dimensional physical environment. The gaming console analyzes the information from the photo-sensors **744** and emitters to determine position and orientation information related to the head-mounted display **102**.

The foregoing components of head-mounted display 102 have been described as merely exemplary components that may be included in head-mounted display 102. In various embodiments of the disclosure, the head-mounted display 102 may or may not include some of the various aforementioned components. Embodiments of the head-mounted display 102 may additionally include other components not presently described, but known in the art, for purposes of facilitating aspects of the present disclosure as herein described.

It will be appreciated by those skilled in the art that in various embodiments of the disclosure, the aforementioned handheld device may be utilized in conjunction with an interactive application displayed on a display to provide various interactive functions. The exemplary embodiments described herein are provided by way of example only, and not by way of limitation.

FIG. 8 is a block diagram of a Game System 800, according to various embodiments of the disclosure. Game System 800 is configured to provide a video stream to one or more Clients 810 via a Network 815. Game System 800 typically includes a Video Server System 820 and an optional game server 825. Video Server System 820 is configured to provide the video stream to the one or more Clients 810 with a minimal quality of service. For example, Video Server System 820 may receive a game command that changes the state of or a point of view within a video game, and provide Clients 810 with an updated video stream reflecting this change in state with minimal lag time. The Video Server System 820 may be configured to provide the video stream in a wide variety of alternative video formats, including formats yet to be defined. Further, the video stream may include video frames configured for presentation to a user at a wide variety of frame rates. Typical frame rates are 30 frames per second, 80 frames per second, and 820 frames per second. Although higher or lower frame rates are included in alternative embodiments of the disclosure.

Clients 810, referred to herein individually as 810A, 810B, etc., may include head mounted displays, terminals, personal computers, game consoles, tablet computers, telephones, set top boxes, kiosks, wireless devices, digital pads, stand-alone devices, handheld game playing devices, and/or the like. Typically, Clients 810 are configured to receive encoded video streams (i.e., compressed), decode the video streams, and present the resulting video to a user, e.g., a player of a game. The processes of receiving encoded video streams and/or decoding the video streams typically includes

storing individual video frames in a receive buffer of the client. The video streams may be presented to the user on a display integral to Client 810 or on a separate device such as a monitor or television. Clients 810 are optionally configured to support more than one game player. For example, a 5 game console may be configured to support two, three, four or more simultaneous players. Each of these players may receive a separate video stream, or a single video stream may include regions of a frame generated specifically for each player, e.g., generated based on each player's point of 10 view. Clients 810 are optionally geographically dispersed. The number of clients included in Game System 800 may vary widely from one or two to thousands, tens of thousands, or more. As used herein, the term "game player" is used to refer to a person that plays a game and the term "game 15 playing device" is used to refer to a device used to play a game. In some embodiments, the game playing device may refer to a plurality of computing devices that cooperate to deliver a game experience to the user. For example, a game console and an HMD may cooperate with the video server 20 system 820 to deliver a game viewed through the HMD. In one embodiment, the game console receives the video stream from the video server system 820, and the game console forwards the video stream, or updates to the video stream, to the HMD for rendering.

Clients **810** are configured to receive video streams via Network **815**. Network **815** may be any type of communication network including, a telephone network, the Internet, wireless networks, powerline networks, local area networks, wide area networks, private networks, and/or the like. In 30 typical embodiments, the video streams are communicated via standard protocols, such as TCP/IP or UDP/IP. Alternatively, the video streams are communicated via proprietary standards

A typical example of Clients **810** is a personal computer 35 comprising a processor, non-volatile memory, a display, decoding logic, network communication capabilities, and input devices. The decoding logic may include hardware, firmware, and/or software stored on a computer readable medium. Systems for decoding (and encoding) video 40 streams are well known in the art and vary depending on the particular encoding scheme used.

Clients 810 may, but are not required to, further include systems configured for modifying received video. For example, a client may be configured to perform further 45 rendering, to overlay one video image on another video image, to crop a video image, and/or the like. For example, Clients 810 may be configured to receive various types of video frames, such as I-frames, P-frames and B-frames, and to process these frames into images for display to a user. In 50 some embodiments, a member of Clients 810 is configured to perform further rendering, shading, conversion to 3-D, or like operations on the video stream. A member of Clients **810** is optionally configured to receive more than one audio or video stream. Input devices of Clients 810 may include, 55 for example, a one-hand game controller, a two-hand game controller, a gesture recognition system, a gaze recognition system, a voice recognition system, a keyboard, a joystick, a pointing device, a force feedback device, a motion and/or location sensing device, a mouse, a touch screen, a neural 60 interface, a camera, input devices yet to be developed, and/or the like.

The video stream (and optionally audio stream) received by Clients **810** is generated and provided by Video Server System **820**. As is described further elsewhere herein, this 65 video stream includes video frames (and the audio stream includes audio frames). The video frames are configured

(e.g., they include pixel information in an appropriate data structure) to contribute meaningfully to the images displayed to the user. As used herein, the term "video frames" is used to refer to frames including predominantly information that is configured to contribute to, e.g. to effect, the images shown to the user. Most of the teachings herein with regard to "video frames" can also be applied to "audio frames."

Clients 810 are typically configured to receive inputs from a user. These inputs may include game commands configured to change the state of the video game or otherwise affect gameplay. The game commands can be received using input devices and/or may be automatically generated by computing instructions executing on Clients 810. The received game commands are communicated from Clients 810 via Network 815 to Video Server System 820 and/or Game Server 825. For example, in some embodiments, the game commands are communicated to Game Server 825 via Video Server System 820. In some embodiments, separate copies of the game commands are communicated from Clients 810 to Game Server 825 and Video Server System **820**. The communication of game commands is optionally dependent on the identity of the command Game commands are optionally communicated from Client 810A through a different route or communication channel that that used to provide audio or video streams to Client 810A.

Game Server 825 is optionally operated by a different entity than Video Server System 820. For example, Game Server 825 may be operated by the publisher of a multiplayer game. In this example, Video Server System 820 is optionally viewed as a client by Game Server 825 and optionally configured to appear from the point of view of Game Server 825 to be a prior art client executing a prior art game engine. Communication between Video Server System 820 and Game Server 825 optionally occurs via Network 815. As such, Game Server 825 can be a prior art multiplayer game server that sends game state information to multiple clients, one of which is game server system 820. Video Server System 820 may be configured to communicate with multiple instances of Game Server 825 at the same time. For example, Video Server System 820 can be configured to provide a plurality of different video games to different users. Each of these different video games may be supported by a different Game Server 825 and/or published by different entities. In some embodiments, several geographically distributed instances of Video Server System 820 are configured to provide game video to a plurality of different users. Each of these instances of Video Server System 820 may be in communication with the same instance of Game Server 825. Communication between Video Server System 820 and one or more Game Server 825 optionally occurs via a dedicated communication channel. For example, Video Server System 820 may be connected to Game Server 825 via a high bandwidth channel that is dedicated to communication between these two systems.

Video Server System 820 comprises at least a Video Source 830, an I/O Device 845, a Processor 850, and non-transitory Storage 855. Video Server System 820 may include one computing device or be distributed among a plurality of computing devices. These computing devices are optionally connected via a communications system such as a local area network.

Video Source 830 is configured to provide a video stream, e.g., streaming video or a series of video frames that form a moving picture. In some embodiments, Video Source 830 includes a video game engine and rendering logic. The video game engine is configured to receive game commands from

a player and to maintain a copy of the state of the video game based on the received commands. This game state includes the position of objects in a game environment, as well as typically a point of view. The game state may also include properties, images, colors and/or textures of objects.

The game state is typically maintained based on game rules, as well as game commands such as move, turn, attack, set focus to, interact, use, and/or the like. Part of the game engine is optionally disposed within Game Server 825. Game Server 825 may maintain a copy of the state of the 10 game based on game commands received from multiple players using geographically disperse clients. In these cases, the game state is provided by Game Server 825 to Video Source 830, wherein a copy of the game state is stored and rendering is performed. Game Server 825 may receive game 15 commands directly from Clients 810 via Network 815, and/or may receive game commands via Video Server System 820.

Video Source 830 typically includes rendering logic, e.g., hardware, firmware, and/or software stored on a computer 20 readable medium such as Storage 855. This rendering logic is configured to create video frames of the video stream based on the game state. All or part of the rendering logic is optionally disposed within a graphics processing unit (GPU). Rendering logic typically includes processing stages 25 configured for determining the three-dimensional spatial relationships between objects and/or for applying appropriate textures, etc., based on the game state and viewpoint. The rendering logic produces raw video that is then usually encoded prior to communication to Clients 810. For 30 example, the raw video may be encoded according to an Adobe Flash® standard, .wav, H.264, H.263, On2, VP6, VC-1, WMA, Huffyuv, Lagarith, MPG-x. Xvid. FFmpeg, x264, VP6-8, realvideo, mp3, or the like. The encoding process produces a video stream that is optionally packaged 35 for delivery to a decoder on a remote device. The video stream is characterized by a frame size and a frame rate. Typical frame sizes include 800×600, 1280×720 (e.g., 720p), 1024×768, although any other frame sizes may be used. The frame rate is the number of video frames per 40 second. A video stream may include different types of video frames. For example, the H.264 standard includes a "P" frame and a "I" frame. I-frames include information to refresh all macro blocks/pixels on a display device, while P-frames include information to refresh a subset thereof. 45 P-frames are typically smaller in data size than are I-frames. As used herein the term "frame size" is meant to refer to a number of pixels within a frame. The term "frame data size" is used to refer to a number of bytes required to store the frame.

In alternative embodiments Video Source 830 includes a video recording device such as a camera. This camera may be used to generate delayed or live video that can be included in the video stream of a computer game. The images and images recorded using a still or video camera. Video Source 830 may also include storage devices configured to store previously recorded video to be included in a video stream. Video Source 830 may also include motion or positioning sensing devices configured to detect motion or 60 position of an object, e.g., person, and logic configured to determine a game state or produce video-based on the detected motion and/or position.

Video Source 830 is optionally configured to provide overlays configured to be placed on other video. For 65 example, these overlays may include a command interface, log in instructions, messages to a game player, images of

other game players, video feeds of other game players (e.g., webcam video). In embodiments of Client 810A including a touch screen interface or a gaze detection interface, the overlay may include a virtual keyboard, joystick, touch pad, and/or the like. In one example of an overlay a player's voice is overlaid on an audio stream. Video Source 830 optionally further includes one or more audio sources.

In embodiments wherein Video Server System 820 is configured to maintain the game state based on input from more than one player, each player may have a different point of view comprising a position and direction of view. Video Source 830 is optionally configured to provide a separate video stream for each player based on their point of view. Further, Video Source 830 may be configured to provide a different frame size, frame data size, and/or encoding to each of Client 810. Video Source 830 is optionally configured to provide 3-D video.

I/O Device **845** is configured for Video Server System **820** to send and/or receive information such as video, commands, requests for information, a game state, gaze information, device motion, device location, user motion, client identities, player identities, game commands, security information, audio, and/or the like. I/O Device 845 typically includes communication hardware such as a network card or modem. I/O Device 845 is configured to communicate with Game Server 825, Network 815, and/or Clients 810.

Processor 850 is configured to execute logic, e.g. software, included within the various components of Video Server System 820 discussed herein. For example, Processor 850 may be programmed with software instructions in order to perform the functions of Video Source 830, Game Server 825, and/or a Client Qualifier 860. Video Server System 820 optionally includes more than one instance of Processor 850. Processor 850 may also be programmed with software instructions in order to execute commands received by Video Server System 820, or to coordinate the operation of the various elements of Game System 800 discussed herein. Processor 850 may include one or more hardware device. Processor 850 is an electronic processor.

Storage 855 includes non-transitory analog and/or digital storage devices. For example, Storage 855 may include an analog storage device configured to store video frames. Storage 855 may include a computer readable digital storage, e.g., a hard drive, an optical drive, or solid state storage. Storage 855 is configured (e.g., by way of an appropriate data structure or file system) to store video frames, artificial frames, a video stream including both video frames and artificial frames, audio frame, an audio stream, and/or the like. Storage 855 is optionally distributed among a plurality of devices. In some embodiments, Storage 855 is configured to store the software components of Video Source 830 discussed elsewhere herein. These components may be stored in a format ready to be provisioned when needed.

Video Server System 820 optionally further comprises resulting video stream optionally includes both rendered 55 Client Qualifier 860. Client Qualifier 860 is configured for remotely determining the capabilities of a client, such as Clients 810A or 810B. These capabilities can include both the capabilities of Client 810A itself as well as the capabilities of one or more communication channels between Client 810A and Video Server System 820. For example, Client Qualifier 860 may be configured to test a communication channel through Network 815.

> Client Qualifier 860 can determine (e.g., discover) the capabilities of Client 810A manually or automatically. Manual determination includes communicating with a user of Client 810A and asking the user to provide capabilities. For example, in some embodiments, Client Qualifier 860 is

configured to display images, text, and/or the like within a browser of Client **810**A. In one embodiment, Client **810**A is an HMD that includes a browser. In another embodiment, client **810**A is a game console having a browser, which may be displayed on the HMD. The displayed objects request that 5 the user enter information such as operating system, processor, video decoder type, type of network connection, display resolution, etc., of Client **810**A. The information entered by the user is communicated back to Client Qualifier

Automatic determination may occur, for example, by execution of an agent on Client **810**A and/or by sending test video to Client **810**A. The agent may comprise computing instructions, such as java script, embedded in a web page or installed as an add-on. The agent is optionally provided by 15 Client Qualifier **860**. In various embodiments, the agent can find out processing power of Client **810**A, decoding and display capabilities of Client **810**A, lag time reliability and bandwidth of communication channels between Client **810**A and Video Server System **820**, a display type of Client **810**A, of firewalls present on Client **810**A, hardware of Client **810**A, software executing on Client **810**A, registry entries within Client **810**A, and/or the like.

Client Qualifier 860 includes hardware, firmware, and/or software stored on a computer readable medium. Client 25 Qualifier 860 is optionally disposed on a computing device separate from one or more other elements of Video Server System 820. For example, in some embodiments, Client Qualifier 860 is configured to determine the characteristics of communication channels between Clients 810 and more 30 than one instance of Video Server System 820. In these embodiments the information discovered by Client Qualifier can be used to determine which instance of Video Server System 820 is best suited for delivery of streaming video to one of Clients 810.

It should be understood that the various embodiments defined herein may be combined or assembled into specific implementations using the various features disclosed herein. Thus, the examples provided are just some possible examples, without limitation to the various implementations 40 that are possible by combining the various elements to define many more implementations. In some examples, some implementations may include fewer elements, without departing from the spirit of the disclosed or equivalent implementations.

Embodiments of the present disclosure may be practiced with various computer system configurations including hand-held devices, microprocessor-based or programmable consumer electronics, minicomputers, mainframe computers and the like. Embodiments of 50 the present disclosure can also be practiced in distributed computing environments where tasks are performed by remote processing devices that are linked through a wire-based or wireless network.

With the above embodiments in mind, it should be understood that embodiments of the present disclosure can employ various computer-implemented operations involving data stored in computer systems. These operations are those requiring physical manipulation of physical quantities. Any of the operations described herein that form part of 60 embodiments of the present disclosure are useful machine operations. Embodiments of the invention also relate to a device or an apparatus for performing these operations. The apparatus can be specially constructed for the required purpose, or the apparatus can be a general-purpose computer 65 selectively activated or configured by a computer program stored in the computer. In particular, various general-pur-

32

pose machines can be used with computer programs written in accordance with the teachings herein, or it may be more convenient to construct a more specialized apparatus to perform the required operations.

The disclosure can also be embodied as computer readable code on a computer readable medium. The computer readable medium is any data storage device that can store data, which can be thereafter be read by a computer system. Examples of the computer readable medium include hard drives, network attached storage (NAS), read-only memory, random-access memory, CD-ROMs, CD-Rs, CD-RWs, magnetic tapes and other optical and non-optical data storage devices. The computer readable medium can include computer readable tangible medium distributed over a network-coupled computer system so that the computer readable code is stored and executed in a distributed fashion.

Although the method operations were described in a specific order, it should be understood that other housekeeping operations may be performed in between operations, or operations may be adjusted so that they occur at slightly different times, or may be distributed in a system which allows the occurrence of the processing operations at various intervals associated with the processing, as long as the processing of the overlay operations are performed in the desired way.

Although the foregoing disclosure has been described in some detail for purposes of clarity of understanding, it will be apparent that certain changes and modifications can be practiced within the scope of the appended claims. Accordingly, the present embodiments are to be considered as illustrative and not restrictive, and embodiments of the present disclosure is not to be limited to the details given herein, but may be modified within the scope and equivalents of the appended claims.

What is claimed is:

1. A method for implementing a graphics pipeline, comprising:

building a first shadow map of high resolution;

building a second shadow map based on the first shadow map, wherein the second shadow map has a lower resolution than the first shadow map;

determining a light source affecting a virtual scene;

projecting geometries of objects of an image of the virtual scene onto a plurality of pixels of a display from a first point-of-view;

determining a first set of geometries is drawn to a first pixel, wherein the first set of geometries includes at least one fragment;

determining the first set of geometries is in shadow based on the light source;

determining a foveal region when rendering the image of the virtual scene, wherein the foveal region corresponds to where an attention of a user is directed using a fragment shader of the graphics pipeline;

determining the first set of geometries is outside of the foveal region using the fragment shader;

rendering the first set of geometries for the first pixel using the second shadow map; and

disabling other shadowing effects during rendering of the first set of geometries for the first pixel.

2. The method of claim 1, further comprising:

determining that a second set of geometries is drawn to a second pixel;

determining the second set of geometries is inside the foveal region; and

rendering the second set of geometries for the second pixel using the first shadow map.

3. The method of claim 1, further comprising:

building a series of shadow maps based on distances from the foveal region, wherein resolutions of shadow maps in the series decrease as corresponding distances from the foveal region increases;

determining a third set of geometries is drawn to a third pixel that is outside the foveal region;

determining that the third set of geometries is in shadow based on the light source;

determining a distance of the third set of geometries as 10 drawn in the third pixel from the foveal region;

selecting a third shadow map from the series of shadow maps based on the distance of the third set of geometries; and

rendering the third set of geometries for the third pixel 15 using the third shadow map.

4. The method of claim **1**, further comprising:

centering the foveal region to a center of the display, wherein the foveal region is static.

5. The method of claim 1, further comprising:

tracking gaze of the user viewing the virtual scene, wherein the gaze indicates a direction into the virtual scene towards which the attention of the user is directed:

determining a first direction of the user viewing the image 25 of the virtual scene; and

centering the foveal region in the image based on the first direction.

6. The method of claim **1**, wherein the display comprises a head mounted display.

7. A computer system comprising:

a processor;

memory coupled to the processor and having stored therein instructions that, if executed by the computer system, cause the computer system to execute a method 35 for implementing a graphics pipeline, comprising: building a first shadow map of high resolution;

building a second shadow map based on the first shadow map, wherein the second shadow map has a lower resolution than the first shadow map;

determining a light source affecting a virtual scene; projecting geometries of objects of an image of the virtual scene onto a plurality of pixels of a display from a first point-of-view;

determining a first set of geometries is drawn to a first 45 pixel, wherein the first set of geometries includes at least one fragment:

determining the first set of geometries is in shadow based on the light source;

determining a foveal region when rendering the image 50 of the virtual scene, wherein the foveal region corresponds to where an attention of a user is directed using a fragment shader of the graphics pipeline;

determining the first set of geometries is outside of the foveal region using the fragment shader;

55

rendering the first set of geometries for the first pixel using the second shadow map; and

disabling other shadowing effects during rendering of the first set of geometries for the first pixel.

8. The computer system of claim **7**, wherein the method 60 further comprises:

determining that a second set of geometries is drawn to a second pixel;

determining the second set of geometries is inside the foveal region; and

rendering the second set of geometries for the second pixel using the first shadow map.

34

9. The computer system of claim **7**, wherein the method further comprises:

building a series of shadow maps based on distances from the foveal region, wherein resolutions of shadow maps in the series decrease as corresponding distances from the foveal region increases;

determining a third set of geometries is drawn to a third pixel that is outside the foveal region;

determining that the third set of geometries is in shadow based on the light source;

determining a distance of the third set of geometries as drawn in the third pixel from the foveal region;

selecting a third shadow map from the series of shadow maps based on the distance of the third set of geometries; and

rendering the third set of geometries for the third pixel using the third shadow map.

10. The computer system of claim 7, wherein the method 20 further comprises:

centering the foveal region to a center of the display, wherein the foveal region is static.

11. The computer system of claim 7, wherein the method further comprises:

tracking gaze of the user viewing the virtual scene, wherein the gaze indicates a direction into the virtual scene towards which the attention of the user is directed:

determining a first direction of the user viewing the image of the virtual scene; and

centering the foveal region in the image based on the first direction.

12. The computer system of claim 7, wherein in the method the display comprises a head mounted display.

13. A non-transitory computer-readable medium storing a computer program for implementing a graphics pipeline, the computer-readable medium comprising:

program instructions for building a first shadow map of high resolution;

program instructions for building a second shadow map based on the first shadow map, wherein the second shadow map has a lower resolution than the first shadow map;

program instructions for determining a light source affecting a virtual scene;

program instructions for projecting geometries of objects of an image of the virtual scene onto a plurality of pixels of a display from a first point-of-view;

program instructions for determining a first set of geometries is drawn to a first pixel wherein the first set of geometries includes at least one fragment;

program instructions for determining the first set of geometries is in shadow based on the light source;

program instructions for determining a foveal region when rendering the image of the virtual scene, wherein the foveal region corresponds to where an attention of a user is directed using a fragment shader of the graphics pipeline;

program instructions for determining the first set of geometries is outside of the foveal region using the fragment shader;

program instructions for rendering the first set of geometries for the first pixel using the second shadow map; and

program instructions for disabling other shadowing effects during rendering of the first set of geometries for the first pixel.

14. The computer-readable medium of claim **13**, further comprising:

program instructions for determining that a second set of geometries is drawn to a second pixel;

program instructions for determining the second set of ⁵ geometries is inside the foveal region; and

program instructions for rendering the second set of geometries for the second pixel using the first shadow map.

15. The computer-readable medium of claim 13, further comprising:

program instructions for building a series of shadow maps based on distances from the foveal region, wherein resolutions of shadow maps in the series decrease as corresponding distances from the foveal region increases:

program instructions for determining a third set of geometries is drawn to a third pixel that is outside the foveal region;

program instructions for determining that the third set of geometries is in shadow based on the light source;

program instructions for determining a distance of the third set of geometries as drawn in the third pixel from the foveal region; 36

program instructions for selecting a third shadow map from the series of shadow maps based on the distance of the third set of geometries; and

program instructions for rendering the third set of geometries for the third pixel using the third shadow map.

16. The computer-readable medium of claim 13, further comprising:

program instructions for centering the foveal region to a center of the display, wherein the foveal region is static.

17. The computer-readable medium of claim 13, further comprising:

program instructions for tracking gaze of the user viewing the virtual scene, wherein the gaze indicates a direction into the virtual scene towards which the attention of the user is directed;

program instructions for determining a first direction of the user viewing the image of the virtual scene; and

program instructions for centering the foveal region in the image based on the first direction.

18. The computer-readable medium of claim 13, wherein 20 the program instructions for projecting geometries further comprises:

program instructions for projecting the geometries of objects of the image onto pixels of a head mounted display.

* * * * *